

# CS 559 Machine Learning

## Lecture 4: Model Selection and Logistic Regression

Ping Wang

Department of Computer Science  
Stevens Institute of Technology



# Today's Lecture

- Model Selection
- Evaluation Metrics for Classification
- Probabilistic Generative Models
- Logistic Regression

# Model Selection

# What is Model Selection?

Given a set of models  $M = \{M_1, M_2, \dots, M_R\}$ , choose the model that is expected to do the best on **the test data**.  $M$  may consist of:

- Same learning model with **different complexities** or **hyperparameters**.
  - Nonlinear regression: polynomials with different degrees
  - K-Nearest Neighbors: Different choices of  $K$
  - Decision Trees: Different choices of the number of levels/leaves
  - SVM: Different choices of the misclassification penalty
  - Regularized models: Different choices of the regularization parameter
  - Kernel based methods: Different choices of kernels ...and almost any learning problem
- Different **learning models** (e.g. SVM, kNN, DT, etc)
- Note: usually considered in supervised learning but unsupervised learning faces this issue too.

# Held-out Data

- Set aside a fraction (10-20%) of the training data.
- This part becomes our held-out data (validation/development)
- Remember: Held-out data is NOT the test data
- Train each model using the remaining training data
- Evaluate error on the held-out data
- Choose the model with the smallest held-out error
- Problems:
  - Wastes training data
  - If there was an unfortunate split, the data distributions may be changed (can be alleviated by repeated random subsampling).

# Cross-Validation

- $K$ -fold Cross-Validation on  $N$  training examples
- Create  $K$  equal sized partitions of the training data
- Each partition has  $N/K$  examples
- Train using  $K - 1$  partitions, validate on the remaining partition
- Repeat the same  $K$  times, each with a different validation partition
- Choose the model with the smallest average validation error
- Usually  $K$  is chosen as 5 or 10.

# Leave-One-Out (LOO) Cross-Validation

- Special case of  $K$ -fold Cross-Validation when  $K = N$
- Each partition is now **an example**
- Train using  $N - 1$  examples, validate on the remaining example
- Repeat the same  $N$  times, each with a different validation example
- Choose the model with the **smallest average validation error**
- **Can be expensive** for large  $N$ . Typically used when  $N$  is small

# Random Subsampling Cross-Validation

- Randomly subsample a fixed fraction  $\alpha N$  ( $0 < \alpha < 1$ ) of examples; call it the validation set
- Training using the rest of the examples, measure error on the validation set
- Repeat  $K$  times, each with a different randomly chosen validation set
- Choose the model with the smallest average validation error
- Usually  $\alpha$  is chosen as 0.1,  $K$  as 10



# Bootstrapping

- Given a set of  $N$  examples
- Idea: Sample  $N$  elements from this set with **replacement** (already sampled elements can be picked again)
- Use this new set as the training data
- The set of examples not selected as the validation data
- For large  $N$ , training data consists of about only **63% unique** examples

- **Expected model error:**

$$e = 0.632e_{test} + 0.368e_{training}$$

- This can break down if we overfit and  $e_{training} = 0$

# Evaluation Metrics for Classification

# Metrics for Performance Evaluation

- Focus on the **predictive capability** of a model
  - Rather than how fast it takes to classify or build models, interpretability, scalability, etc.
- The **contingency table** or **confusion matrix**:

Predicted Class	Actual Class		
		Positive	Negative
	Positive	True positive (TP)	False positive (FP)
	Negative	False negative (FN)	True negative (TN)

# Metrics for Performance Evaluation

Predicted Class	Actual Class		
		Positive	Negative
	Positive	True positive (TP)	False positive (FP)
	Negative	False negative (FN)	True negative (TN)

- Most widely-used metric: **Accuracy** =  $\frac{TP+TN}{TP+TN+FP+FN}$
- Limitation:
  - For a 2-class problem,  $N_0 = 9990, N_1 = 10$ .
  - If the model predicts everything to be  $C_0$ , the accuracy is  $\frac{9990}{10000} = 99.9\%$ .
  - This is misleading since the model does not detect any example from  $C_1$ .

# Cost of Classification

- Consider the **cost of misclassification**
- Example: cancer vs. non-cancer

Cost Matrix	Actual Class		
		Positive	Negative
	Predicted Class		
	Positive	-1	1
	Negative	100	0

Model 1	Actual Class		
Predicted Class		Positive	Negative
	Positive	150	60
	Negative	40	250

Accuracy=80%  
Cost=3910

Model 2	Actual Class		
Predicted Class		Positive	Negative
	Positive	250	5
	Negative	45	200

Accuracy=90%  
Cost=4255

# Cost-Sensitive Classification

- So far, we haven't considered costs during training.
- Most learning algorithms do not perform cost-sensitive learning.
- Taking the costs into the training procedure will be an **algorithm specific task**.
  - Consider the cost when making the predictions and minimize the expected cost.
  - Only predict high-cost class when the model is very confident about the prediction.

# Metrics for Performance Evaluation

Predicted Class	Actual Class		
		Positive	Negative
	Positive	True positive (TP)	False positive (FP)
	Negative	False negative (FN)	True negative (TN)

- **Recall** =  $\frac{TP}{TP+FN}$ , the fraction of relevant instances that were retrieved.
- **Precision** =  $\frac{TP}{TP+FP}$ , the fraction of relevant instances among the retrieved instances.
- **F-measure** is the harmonic mean between recall and precision.

$$F - measure = \frac{2Recall \times Precision}{Recall + Precision} = \frac{2TP}{2TP + FP + FN}$$

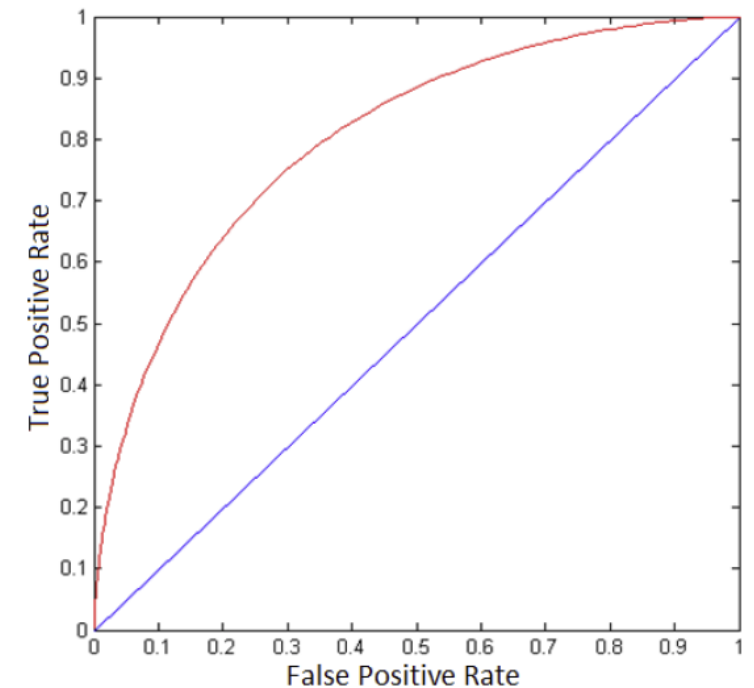
# ROC Curve

- An ROC curve (receiver operating characteristic curve) is a graph showing the performance of a classification model at all classification thresholds.
- This curve plots two parameters:
  - True positive rate on y-axis:  $TPR = \frac{TP}{TP+FN}$
  - False positive rate on x-axis:  $FPR = \frac{FP}{FP+TN}$
- Performance of each classifier represented as a point on the ROC curve.
  - Changing the threshold of algorithm changes the location of the point.
  - Lowering the classification threshold classifies more items as positive, thus increasing both False Positives and True Positives.



# ROC Curve

- (TPR, FPR)
  - (0,0): declare everything to be negative class
  - (1,1): declare everything to be positive class
  - (1,0): Ideal/perfect classification
- Diagonal line represents random guessing. Points above the diagonal represent good classification results.
- **AUC: Area Under the ROC Curve**
  - Measures the entire two-dimensional area under the ROC curve
  - Provides an aggregate measure of performance across all possible classification thresholds.
  - Ideal:  $AUC=1$ ; Random guess:  $AUC=0.5$



# How to Construct an ROC Curve

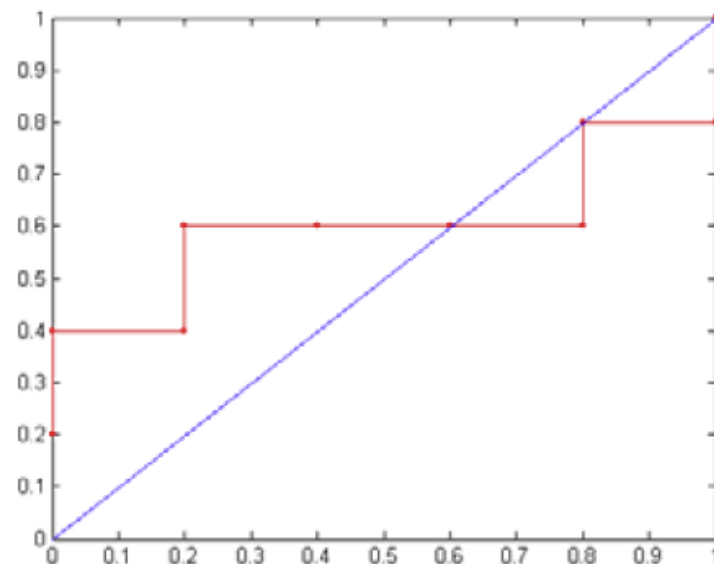
1. Use a classifier to produce posterior probability for each test instance  $P(+|A)$
2. Sort the instances according to  $P(+|A)$  in **decreasing order**
3. Apply threshold at **each unique value** of  $P(+|A)$
4. Count the number of TP, FP, TN, FN at each threshold
5.  $TPR = TP/(TP+FN)$ ;  $FPR = FP/(FP + TN)$

Instance	$P(+ A)$	True Class
1	0.95	+
2	0.93	+
3	0.87	-
4	0.85	-
5	0.85	-
6	0.85	+
7	0.76	-
8	0.53	+
9	0.43	-
10	0.25	+

# How to Construct an ROC Curve

Class	+	-	+	-	-	-	+	-	+	+	
Threshold $\geq$	0.25	0.43	0.53	0.76	0.85	0.85	0.85	0.87	0.93	0.95	1.00
TP	5	4	4	3	3	3	3	2	2	1	0
FP	5	5	4	4	3	2	1	1	0	0	0
TN	0	0	1	1	2	3	4	4	5	5	5
FN	0	1	1	2	2	2	2	3	3	4	5
TPR	1	0.8	0.8	0.6	0.6	0.6	0.6	0.4	0.4	0.2	0
FPR	1	1	0.8	0.8	0.6	0.4	0.2	0.2	0	0	0

ROC Curve



Instance	P(+ A)	True Class
1	0.95	+
2	0.93	+
3	0.87	-
4	0.85	-
5	0.85	-
6	0.85	+
7	0.76	-
8	0.53	+
9	0.43	-
10	0.25	+

# Feature Selection

# Feature Selection

Selecting a useful subset from all the features. Why?

- Some algorithms **scale (computationally) poorly** with increased dimension
- **Irrelevant** features can confuse some algorithms
- **Redundant** features adversely affect regularization
- Removal of features can **increase (relative) margin** (and generalization)
- Reduces data set and resulting model size
- Note: Feature Selection is different from Feature Extraction
  - The latter transforms original features to get a small set of new features (we will discuss in dimensionality reduction).

# Feature Selection Methods

- Methods agnostic to the learning algorithm
  - Preprocessing based methods
    - E.g., remove a binary feature if it's ON in very few or most examples
  - Filter Feature Selection methods
    - Use some ranking criteria to rank features
    - Select the top-ranking features
- Wrapper Methods (keep the learning algorithm in the loop)
  - Forward/ Backward Search
  - Requires repeated runs of the learning algorithm with different set of features
  - Can be computationally expensive

# Filter Feature Selection

- Uses heuristics but is much faster than wrapper methods
- **Correlation Criteria**: Rank features based on their correlation with the labels

$$R(X_d, y) = \frac{cov(X_d, y)}{\sqrt{var(X_d)var(y)}}$$

- **Mutual Information Criteria**

$$MI(X_d, y) = \sum_{X_d \in \{0,1\}} \sum_{y \in \{-1,1\}} P(X_d, y) \log \frac{P(X_d, y)}{P(X_d)P(y)}$$

- High mutual information means high relevance of that feature
- These probabilities can be easily estimated from the data

# Wrapper Methods

- Forward Search
  - Start with no features
  - Greedily include the most relevant feature
  - Stop when selected the desired number of features
- Backward Search
  - Start with all features
  - Greedily remove the least relevant feature
  - Stop when selected the desired number of features
- Inclusion/Removal criteria uses cross-validation



# Probabilistic Generative Models

# Probabilistic Generative Models

- We now turn to a **probabilistic** approach to classification.
- Show how models with **linear decision boundaries** arise from simple assumptions about the distribution of the data.
- Adopt the **generative approach**: Solve the inference problem of estimating the class-conditional densities  $p(x|C_k)$  for each class  $C_k$ .
- Infer the prior class probabilities  $p(C_k)$ .
- Use Bayes' theorem to find the class posterior probabilities:

$$p(C_k|x) = \frac{p(x|C_k)p(C_k)}{p(x)}$$

where  $p(x) = \sum_k p(x|C_k)p(C_k)$

# Probabilistic Generative Approach: Two Classes

- Two-class case:

$$p(C_1|x) = \frac{p(x|C_1)p(C_1)}{p(x|C_1)p(C_1) + p(x|C_2)p(C_2)} = \frac{1}{1 + e^{-a}} = \sigma(a)$$

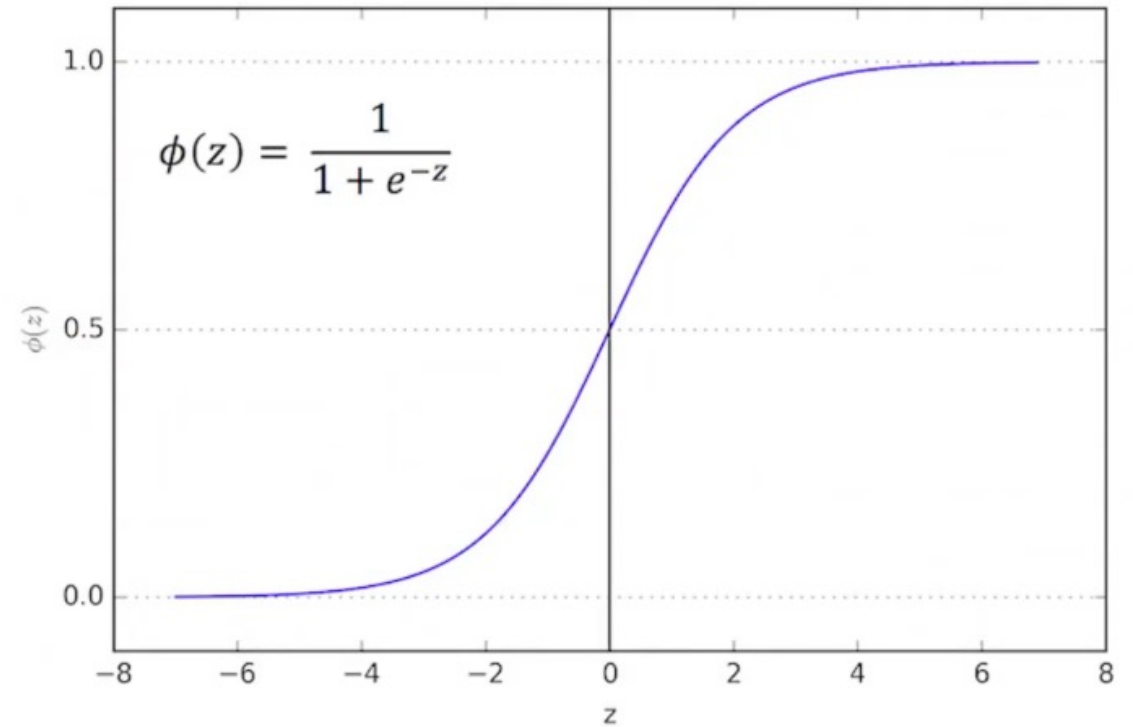
where

$$a = \ln \frac{p(x|C_1)p(C_1)}{p(x|C_2)p(C_2)}$$

$$\sigma(a) = \frac{1}{1+e^{-a}} \text{ (Logistic Sigmoid Function)}$$

# Logistic Sigmoid Function

- It tends to 1 as  $z \rightarrow \infty$
- It tends to 0 as  $z \rightarrow -\infty$
- It is bounded between 0 and 1.



# Probabilistic Generative Approach: Multiple Classes

- Multiple ( $K > 2$ ) classes case:

$$p(C_k|x) = \frac{p(x|C_k)p(C_k)}{\sum_{j=1}^K p(x|C_j)p(C_j)} = \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}}$$

where

$$a_k = \ln p(x|C_k)p(C_k)$$

$$\sigma(a) = \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}} \text{ (softmax function)}$$

# Probabilistic Generative Models

- Let's assume the **class-conditional densities** are **Gaussian** with the same covariance matrix:

$$p(x|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1} (x-\mu_k)}$$

- Two-class case: we can show the following result:

$$\begin{aligned} p(C_1|x) &= \sigma(w^T x + \omega_0) \\ w &= \Sigma^{-1}(\mu_1 - \mu_2) \\ \omega_0 &= -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \ln \frac{p(C_1)}{p(C_2)} \end{aligned}$$

- How to get the result?

# Probabilistic Generative Models

$$p(C_1|x) = \sigma(a) = \frac{1}{1 + e^{-a}} \quad p(x|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1}(x-\mu_k)}$$

$$a = \ln \frac{p(x|C_1)p(C_1)}{p(x|C_2)p(C_2)}$$

$$= \ln p(x|C_1) - \ln p(x|C_2) + \ln \frac{p(C_1)}{p(C_2)} \quad (\text{Replace } p(x|C_k))$$

$$= -\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) + \frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2) + \ln \frac{p(C_1)}{p(C_2)}$$

# Probabilistic Generative Models

$$\begin{aligned} a &= -\frac{1}{2}(x - \mu_1)^T \Sigma^{-1}(x - \mu_1) + \frac{1}{2}(x - \mu_2)^T \Sigma^{-1}(x - \mu_2) + \ln \frac{p(C_1)}{p(C_2)} \\ &= -\frac{1}{2}x^T \Sigma^{-1}x + \frac{1}{2}\mu_1^T \Sigma^{-1}x + \frac{1}{2}x^T \Sigma^{-1}\mu_1 - \frac{1}{2}\mu_1^T \Sigma^{-1}\mu_1 \\ &\quad + \frac{1}{2}x^T \Sigma^{-1}x - \frac{1}{2}\mu_2^T \Sigma^{-1}x - \frac{1}{2}x^T \Sigma^{-1}\mu_2 + \frac{1}{2}\mu_2^T \Sigma^{-1}\mu_2 + \ln \frac{p(C_1)}{p(C_2)} \\ &= \mu_1^T \Sigma^{-1}x - \mu_2^T \Sigma^{-1}x - \frac{1}{2}\mu_1^T \Sigma^{-1}\mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1}\mu_2 + \ln \frac{p(C_1)}{p(C_2)} \\ &= \underbrace{(\mu_1 - \mu_2)^T \Sigma^{-1}}_{w^T} x \underbrace{\left[ -\frac{1}{2}\mu_1^T \Sigma^{-1}\mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1}\mu_2 + \ln \frac{p(C_1)}{p(C_2)} \right]}_{\omega_0} \end{aligned}$$

The quadratic terms in  $x$  have cancelled, leading to a linear function of  $x$ .



# Probabilistic Generative Models

- We have shown:

$$p(C_1|x) = \sigma(w^T x + \omega_0)$$

$$w = \Sigma^{-1}(\mu_1 - \mu_2)$$

$$\omega_0 = -\frac{1}{2}\mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2}\mu_2^T \Sigma^{-1} \mu_2 + \ln \frac{p(C_1)}{p(C_2)}$$

- Decision boundary is **linear in input space**:

$$p(C_1|x) = p(C_2|x) = 0.5$$

$$\Rightarrow \frac{1}{1+e^{-(w^T x + \omega_0)}} = 0.5 \Rightarrow w^T x + \omega_0 = 0$$

# Probabilistic Generative Models

- Multiple ( $K > 2$ ) classes case:

$$p(C_k|x) = \frac{p(x|C_k)p(C_k)}{\sum_{j=1}^K p(x|C_j)p(C_j)} = \frac{e^{a_k}}{\sum_{j=1}^K e^{a_j}}$$

$$a_k = \ln p(x|C_k)p(C_k)$$

- Decision boundary will occur when two of the posterior probabilities are equal, which will be defined by **linear functions of  $x$** :

$$p(C_k|x) = \frac{e^{w_k^T x + \omega_{k0}}}{\sum_j e^{w_j^T x + \omega_{j0}}}$$

$$w_k = \Sigma^{-1} \mu_k$$

$$\omega_{k0} = -\frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \ln p(C_k)$$

# Maximum Likelihood Solution

- Now we have a parametric functional form for the class-conditional densities:

$$p(x|C_k) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} e^{-\frac{1}{2}(x-\mu_k)^T \Sigma^{-1} (x-\mu_k)}$$

- We can estimate the parameters and the prior class probabilities using maximum likelihood.
  - The case of two classes, each having a Gaussian class-conditional density with a shared covariance matrix.
  - Training data:  $\{x_n, y_n\}, n = 1, \dots, N$ 
    - $y_n = 1$  denotes class  $C_1$  ( $N_1$  data samples);
    - $y_n = 0$  denotes class  $C_2$  ( $N_2$  data samples);
    - Prior class probability:  $p(C_1) = \gamma, p(C_2) = 1 - \gamma$

# Maximum Likelihood Solution

- For a data point  $x_n$  from class  $C_1$ , we have  $y_n = 1$  and therefore:

$$p(x_n, C_1) = p(x|C_1)p(C_1) = \gamma \mathcal{N}(x_n|\mu_1, \Sigma)$$

- For a data point  $x_n$  from class  $C_2$ , we have  $y_n = 0$  and therefore:

$$p(x_n, C_2) = p(x|C_2)p(C_2) = (1 - \gamma) \mathcal{N}(x_n|\mu_2, \Sigma)$$

- Assuming observations are drawn independently, the likelihood function is given as follows, where  $y = (y_1, \dots, y_N)^T$ .

$$\begin{aligned} p(y|\gamma, \mu_1, \mu_2, \Sigma) &= \prod_{n=1}^N [p(x_n, C_1)]^{y_n} [p(x_n, C_2)]^{1-y_n} \\ &= \prod_{n=1}^N [\gamma \mathcal{N}(x_n|\mu_1, \Sigma)]^{y_n} [(1 - \gamma) \mathcal{N}(x_n|\mu_2, \Sigma)]^{1-y_n} \end{aligned}$$

# Maximum Likelihood Solution

- We want to find the values of the parameters that **maximize the likelihood function**, i.e., fit a model that best describes the observed data.

$$p(y|\gamma, \mu_1, \mu_2, \Sigma) = \prod_{n=1}^N [\gamma \mathcal{N}(x_n|\mu_1, \Sigma)]^{y_n} [(1 - \gamma) \mathcal{N}(x_n|\mu_2, \Sigma)]^{1-y_n}$$

- As usual, we consider the log of the likelihood:

$$\begin{aligned} \ln p(y|\gamma, \mu_1, \mu_2, \Sigma) = & \sum_{n=1}^N [y_n \ln \gamma + y_n \ln \mathcal{N}(x_n|\mu_1, \Sigma) \\ & + (1 - y_n) \ln(1 - \gamma) + (1 - y_n) \ln \mathcal{N}(x_n|\mu_2, \Sigma)] \end{aligned}$$

# Maximum Likelihood Solution: Parameter $\gamma$

$$\ln p(y|\gamma, \mu_1, \mu_2, \Sigma) = \sum_{n=1}^N [y_n \ln \gamma + y_n \ln \mathcal{N}(x_n|\mu_1, \Sigma) + (1 - y_n) \ln(1 - \gamma) + (1 - y_n) \ln \mathcal{N}(x_n|\mu_2, \Sigma)]$$

- We first maximize the log-likelihood with respect to  $\gamma$  (set derivative to 0).

$$\gamma = \frac{1}{N} \sum_{n=1}^N y_n = \frac{N_1}{N} = \frac{N_1}{N_1 + N_2}$$

- The maximum likelihood estimate of  $\gamma$  is the fraction of points in class  $C_1$ .
- For multi-class: ML estimate for  $p(C_k)$  is given by the fraction of points in the training set in  $C_k$ .

# Maximum Likelihood Solution: Parameter $\mu$

$$\ln p(D|\gamma, \mu_1, \mu_2, \Sigma) = \sum_{n=1}^N [y_n \ln \gamma + y_n \ln \mathcal{N}(x_n|\mu_1, \Sigma) \\ + (1 - y_n) \ln(1 - \gamma) + (1 - y_n) \ln \mathcal{N}(x_n|\mu_2, \Sigma)]$$

- We then maximize the log-likelihood with respect to  $\mu_1$  (set derivative to 0).

$$\mu_1 = \frac{1}{N_1} \sum_{n=1}^N y_n x_n$$

- The maximum likelihood estimate of  $\mu_1$  is the sample mean of all inputs  $x_n$  in class  $C_1$ .
- Similarly, the maximum likelihood estimate of  $\mu_2$  is given by

$$\mu_2 = \frac{1}{N_2} \sum_{n=1}^N (1 - y_n) x_n$$

# Maximum Likelihood Solution: Parameter $\Sigma$

- Maximize the log-likelihood with respect to the covariance matrix  $\Sigma$  (set derivative to 0), we obtain the estimate  $\Sigma_{ML}$

$$\Sigma_{ML} = \frac{N_1}{N} S_1 + \frac{N_2}{N} S_2$$

where

$$S_1 = \frac{1}{N_1} \sum_{x_n \in C_1} (x_n - \mu_1)(x_n - \mu_1)^T \quad S_2 = \frac{1}{N_2} \sum_{x_n \in C_2} (x_n - \mu_2)(x_n - \mu_2)^T$$

- The maximum likelihood estimate of the covariance is given by the weighted average of the sample covariance matrices associated with each of the classes.
- The results extend to  $K$  classes.



# Summary So Far

- We assumed  $p(x|y = 1) \sim \mathcal{N}(\mu_1, \Sigma)$  and  $p(x|y = 0) \sim \mathcal{N}(\mu_2, \Sigma)$ , and two class-probabilities  $p(y = 1)$  and  $p(y = 0)$ .
- This is called a **generative model**, as we have written down a full joint model over the data.
- Violations of the model assumption can lead to “bad” decision boundaries.

# Probabilistic Generative Models: Parameters

- How many parameters did we estimate to fit Gaussian class-conditional densities (the generative approach)?
- Suppose  $d$  is the dimension of the input space.

$$p(C_1) \Rightarrow 1$$

$$2 \text{ mean vectors} \Rightarrow 2d$$

$$\Sigma \Rightarrow d + \frac{d^2 - d}{2} = \frac{d^2 + d}{2}$$

$$Total = 1 + 2d + \frac{d^2 + d}{2} = O(d^2)$$

# Logistic Regression

# Logistic Regression

$$P(C_1|x) = \sigma(w^T x + \omega_0) = f(x)$$

- We use maximum likelihood to determine the parameters of the logistic regression model.

$$\{x_n, y_n\}, n = 1, \dots, N$$

$y_n = 1$  denotes class  $C_1$ ;  $y_n = 0$  denotes class  $C_2$

- We want to find the values of  $w$  that maximize the posterior probabilities associated to the observed data.
- Likelihood function:

$$\begin{aligned} L(w) &= \prod_{n=1}^N p(C_1|x_n)^{y_n} (1 - p(C_1|x_n))^{1-y_n} \\ &= \prod_{n=1}^N f(x_n)^{y_n} (1 - f(x_n))^{1-y_n} \end{aligned}$$

# Logistic Regression

$$P(C_1|x) = \sigma(w^T x + \omega_0) = f(x)$$

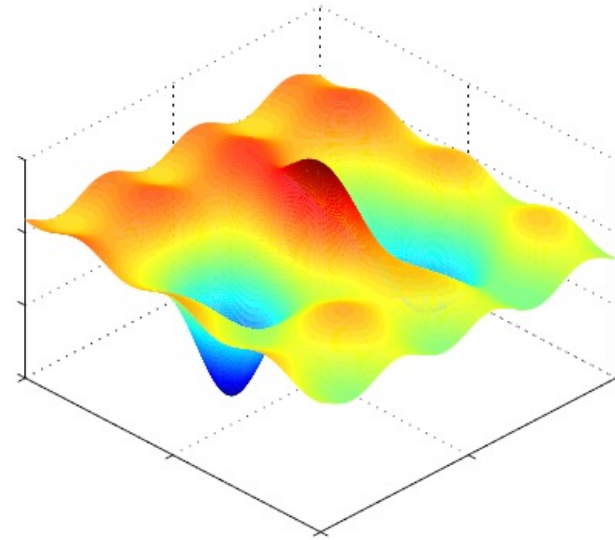
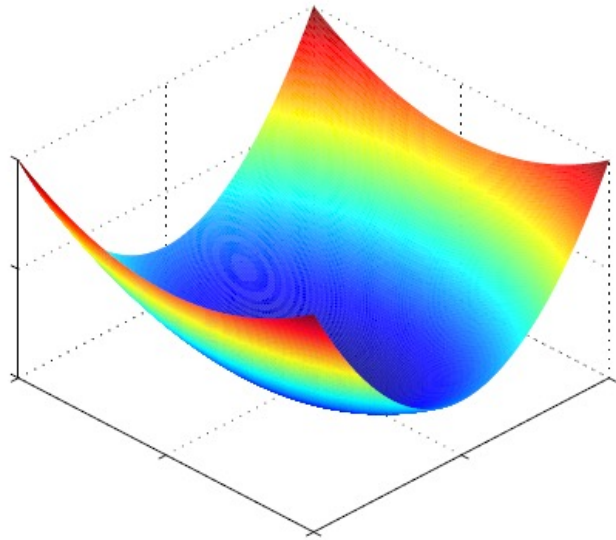
- We consider the **negative logarithm of the likelihood (Cross Entropy)**:

$$\begin{aligned}\varepsilon(w) &= -\ln L(w) = -\ln \prod_{n=1}^N p(C_1|x_n)^{y_n} (1 - p(C_1|x_n))^{1-y_n} \\ &= -\sum_{n=1}^N (y_n \ln f(x_n) + (1 - y_n) \ln(1 - f(x_n)))\end{aligned}$$

- Thus:

$$\max L(w) = \min \varepsilon(w)$$

# The Cost-function for Logistic Regression is Convex.



- Fact: The negative log-likelihood is convex - this makes life much easier.
- There are no local minima to get stuck in, and there is good optimization techniques for convex problems.

# Review of Gradient Descent Algorithms

- Three important elements: data( $x$ ;  $y$ ), loss function, model parameters  $w$
- One important line, gradient update:

$$w^{(t+1)} \leftarrow w^{(t)} - \eta \nabla_w$$

# Gradient Descent

---

```
1: procedure BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:      $g^{(i)}(w) = \text{evaluate\_gradient}(\text{TrainLoss}, \text{data}, w)$ 
4:      $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```

---

## Some Limitations:

- Each iteration requires going over **all training examples (Batch gradient descent)** – expensive and slow when have lots of data.
- Intractable for datasets that don't fit in memory.
- Guaranteed to converge to the global minimum for convex functions, but may end up at a local minimum for non-convex functions.



# Stochastic Gradient Descent (SGD)

---

```
1: procedure STOCHASTIC GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for example  $\in$  data do
5:        $g^{(i)}(w) = \text{evaluate\_gradient}(\text{loss}, \text{example}, w)$ 
6:        $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```

---

- **Shuffling**: to ensure that each data point creates an “independent” change on the model, without being biased by the same points before them.
- Easy to fit in memory as only one data point needs to be processed at a time, thus, computationally less expensive.
- With a high variance that cause the objective function to fluctuate heavily.
- May never reach local minima and oscillate around it due to the fluctuations in each step.

# Mini-batch Gradient Descent

---

```
1: procedure MINI-BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for batch  $\in$  get_batches(data, batch_size=50) do
5:        $g^{(i)}(w) = \text{evaluate\_gradient}(\text{loss}, \text{batch}, w)$ 
6:        $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```

---

- Instead of using the whole data for calculating gradient, we use only a mini-batch of it (batch size is a hyperparameter).
- Reduces the variance of the parameter updates, which can lead to more stable convergence.
- Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient w.r.t. a mini-batch very efficient.

# Logistic Regression: Gradient

$$P(C_1|x) = \sigma(w^T x + \omega_0) = f(x)$$

- We compute the derivative of the error function with respect to  $w$ .

$$\frac{\partial \varepsilon(w)}{\partial w} = \frac{\partial}{\partial w} \left[ -\ln \prod_{n=1}^N p(C_1|x_n)^{y_n} (1 - p(C_1|x_n))^{1-y_n} \right]$$

- The derivative of the logistic sigmoid function:

$$\begin{aligned} \frac{\partial}{\partial a} \sigma(a) &= \frac{\partial}{\partial a} \frac{1}{1 + e^{-a}} = \frac{e^{-a}}{(1 + e^{-a})^2} = \frac{1}{1 + e^{-a}} \frac{e^{-a}}{(1 + e^{-a})} \\ &= \frac{1}{1 + e^{-a}} \left( 1 - \frac{1}{1 + e^{-a}} \right) = \sigma(a)(1 - \sigma(a)) \end{aligned}$$

# Logistic Regression

- The derivative of the error function with respect to  $w$  is:

$$\nabla_w \varepsilon(w) = \sum_{n=1}^N (f(x_n) - y_n) x_n$$

**Homework:** show the detailed steps to calculate this derivation.

- Then, the parameter can be updated by:

$$w^{t+1} := w^t - \eta \nabla_w \varepsilon(w)$$

# Probabilistic Discriminative Models: parameters

- Two-class case:

$$P(C_1|x) = \sigma(w^T x + \omega_0) = f(x)$$

$$P(C_2|x) = 1 - P(C_1|x)$$

- This model is known as **Logistic Regression**.
- Assuming  $x \in \mathbb{R}^d$ , how many parameters do we need to estimate?  
 $d + 1$

# Summary So Far

- From the previous introduction, we know that

$$P(y = 1|x) = \sigma(a(x))$$

where  $\sigma(a) = 1/(1 + \exp(-a))$  and  $a(x) = w^T x + \omega_0$ .

- Notation is simpler if we use 0 and 1 as class labels, so we define  $y_n = 1$  as the label for the positive class, and  $y_n = 0$  a label for the negative class.
- The parameters of  $a(x) = w^T x + \omega_0$  can be learnt by minimizing the negative log-likelihood:

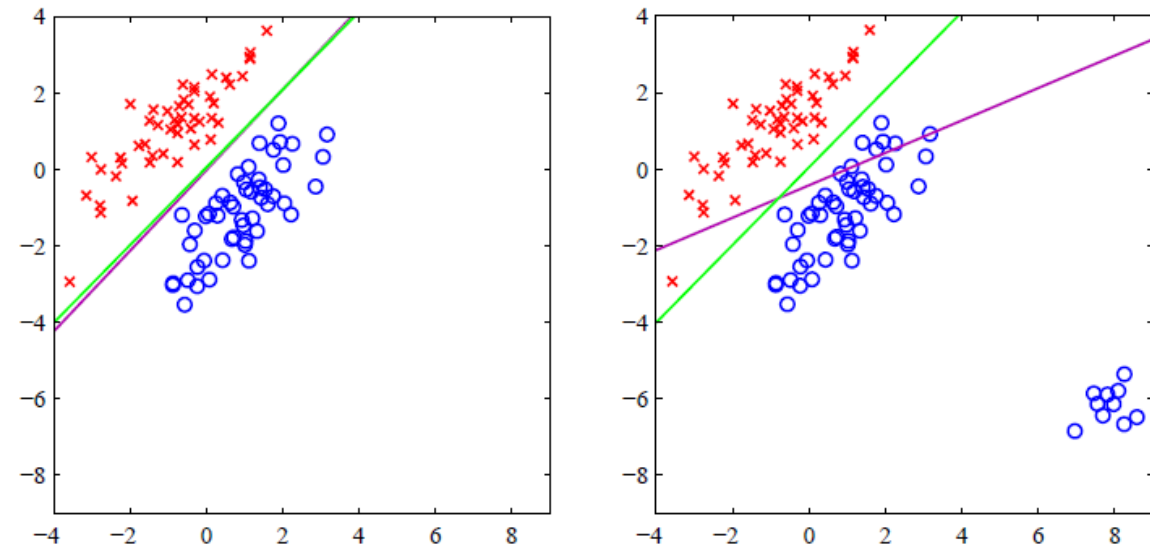
$$L(w) = - \sum_n \{y_n \log p(y_n|x_n, w) + (1 - y_n) \log(1 - p(y_n|x_n, w))\}$$

# Summary So Far

- This is a **discriminative** approach to classification, as we only model the labels, and not the inputs.
- Decision rule and function shape of  $p(y|x)$  will be the same for the generative and the discriminative model, but the parameters were obtained differently.

# Maximum Likelihood Estimation of Logistic Regression

- Logistic regression is a much better algorithm than the algorithms we discussed last week.
- Need to optimize log-likelihood numerically.
- People typically minimize the negative log-likelihood  $\mathcal{L}$  rather than maximizing the log-likelihood.
- To numerically minimize the negative log-likelihood, we need its gradient.



Figures from Bishop PRML, 44a and b



# Multiclass Logistic Regression

- Multiclass case:

$$p(C_k|x) = \frac{e^{w_k^T x + \omega_{k0}}}{\sum_j e^{w_j^T x + \omega_{j0}}} = f_k(x)$$

- We use maximum likelihood to determine the parameters:

$$\{x_n, y_n\}, n = 1, \dots, N$$

$y_n = (0, \dots, 1, \dots, 0)$  denotes class  $C_k$

- We want to find the values of  $w_1, \dots, w_k$  that maximize the posterior probabilities associated to the observed data likelihood function:

$$L(w_1, \dots, w_k) = \prod_{n=1}^N \prod_{k=1}^K p(C_k|x_n)^{y_{nk}} = \prod_{n=1}^N \prod_{k=1}^K f_k(x_n)^{y_{nk}}$$

# Multiclass Logistic Regression

$$L(w_1, \dots, w_k) = \prod_{n=1}^N \prod_{k=1}^K p(C_k | x_n)^{y_{nk}} = \prod_{n=1}^N \prod_{k=1}^K f_k(x_n)^{y_{nk}}$$

- Consider the negative logarithm of the likelihood (cross-entropy error):

$$\varepsilon(w_1, \dots, w_k) = -\ln L(w_1, \dots, w_k) = -\sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(x_n)$$
$$\min_{w_j} \varepsilon((w_j))$$

- The gradient of the error function w.r.t one of the parameter vectors:

$$\frac{\partial}{\partial w_j} \varepsilon(w_1, \dots, w_k) = \frac{\partial}{\partial w_j} \left[ -\sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(x_n) \right]$$

# Multiclass Logistic Regression

$$\frac{\partial}{\partial w_j} \varepsilon(w_1, \dots, w_k) = \frac{\partial}{\partial w_j} \left[ - \sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln f_k(x_n) \right]$$

- The derivatives of the softmax function:

$$\frac{\partial}{\partial a_k} f_k = \frac{\partial}{\partial a_k} \frac{e^{a_k}}{\sum_j e^{a_j}} = \frac{e^{a_k} \sum_j e^{a_j} - e^{a_k} e^{a_k}}{(\sum_j e^{a_j})^2} = f_k - f_k^2 = f_k(1 - f_k)$$

- Thus, for  $j \neq k$

$$\frac{\partial}{\partial a_j} f_k = \frac{\partial}{\partial a_j} \frac{e^{a_k}}{\sum_j e^{a_j}} = \frac{-e^{a_k} e^{a_j}}{(\sum_j e^{a_j})^2} = -f_k f_j$$

- Compact expression ( $I_{kj}$  are the elements of the identity matrix)

$$\frac{\partial}{\partial a_j} f_k = f_k(I_{kj} - f_j)$$

# Multiclass Logistic Regression

$$\nabla_{w_j} \varepsilon(w_1, \dots, w_k) = \sum_{n=1}^N (f_{nj} - y_{nj}) x_n$$

- It can be shown that  $\varepsilon$  is a convex function of  $w$ . Thus, it has a unique minimum.
- Online solution (SGD):

$$w_j^{t+1} = w_j^t - \eta \nabla_{w_j} \varepsilon_n(w) = w_j^t - \eta (f_{nj} - y_{nj}) x_n$$

# Summary of Today's Lecture

- Model Selection
- Evaluation Metrics for Classification
- Probabilistic Generative Models
- Logistic Regression