

CS 559 Machine Learning

Lecture 7: Decision Trees and Ensemble Methods

Ping Wang

Department of Computer Science

Stevens Institute of Technology



Midterm Exam

- Time: 3:30-6:00 PM on March 9, 2023
- Location: Howe 102
- Closed-book exam; one A4 cheatsheet is allowed.
- Calculator is allowed, but phones, laptops, and other devices are not allowed.
- Work on the exam independently.
- You cannot share the calculator or cheatsheet during the exam.

Question Types

- Short answer questions
- Multiple choices questions
- Calculation or equation derivation
- Concept matching for extra points

Question Coverage

Question coverage: Lecture 1-6

1. Perform data preprocessing with feature selection or feature scaling?
2. Maximum Likelihood Estimation: how to estimate the parameters given a probability density function?
3. How to optimize linear regression model using gradient descent/SGD/mini-batch gradient descent method?
4. What is generative/discriminative classification? How to perform classification using Fisher's Linear Discriminant and Perceptron algorithm?
5. Logistic regression: how to perform classification with logistic regression? What is the objective function and how to optimize the model parameters with the gradient descent method?
6. SVM: how to find the maximum margin of the decision boundary? How to transform the data in a nonlinear problem into a new space and solve it using SVM?
7. Model selection and evaluation: such as cross validation, held-out data, and several others. How to evaluate the linear regression methods and classification methods?

Today's Lecture

- Decision Tree
- Ensemble Methods
 - Bagging
 - Boosting
 - Random Forests


Example of a Decision Tree

- Input: 10 data points with three features/attributes
- Goal: predict if a person will be a defaulted borrower.
- Need to find: $f: X \rightarrow Y$

					categorical	categorical	continuous	class
ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower				
1	Yes	Single	125K	No				
2	No	Married	100K	No				
3	No	Single	70K	No				
4	Yes	Married	120K	No				
5	No	Divorced	95K	Yes				
6	No	Married	60K	No				
7	Yes	Divorced	220K	No				
8	No	Single	85K	Yes				
9	No	Married	75K	No				
10	No	Single	90K	Yes				

Decision Trees $f: X \rightarrow Y$

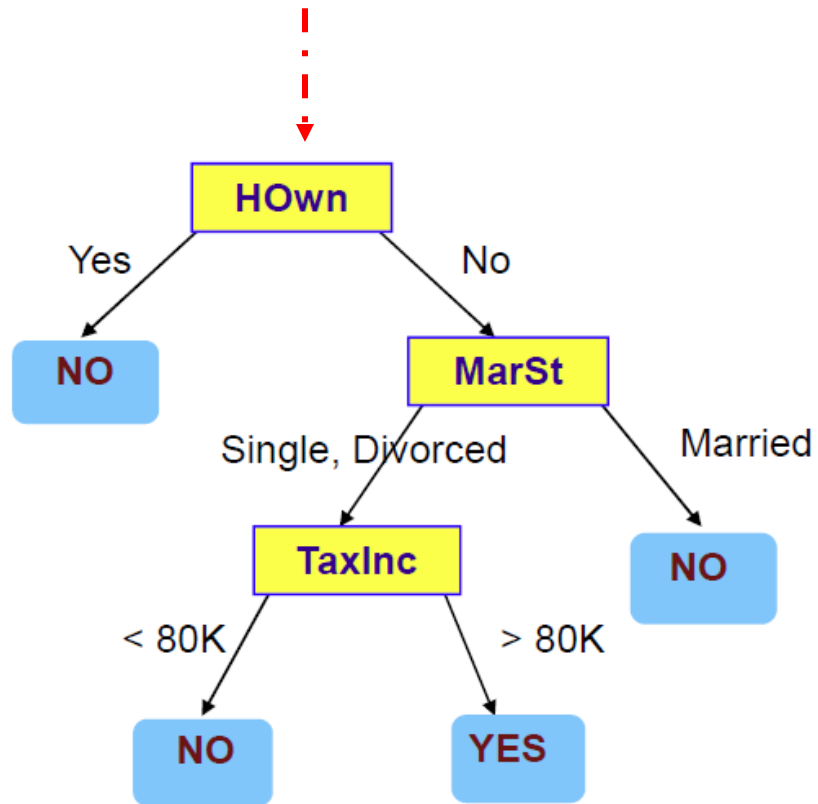
- Each **internal node** tests an attribute x_i
- Each **branch** assigns an attribute value $x_i = v$
- Each **leaf** assigns a class
- To classify input x : traverse the tree from root to leaf, output the labeled y



There could be more than one tree that fits the same data!

Apply Model to Test Data

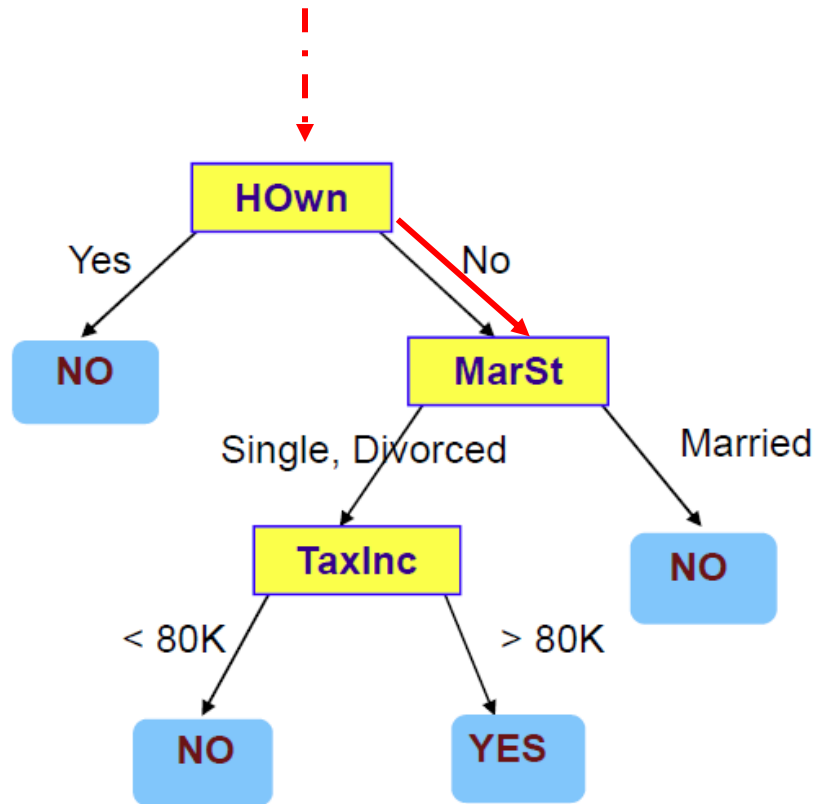
Start from the root of tree



Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Apply Model to Test Data

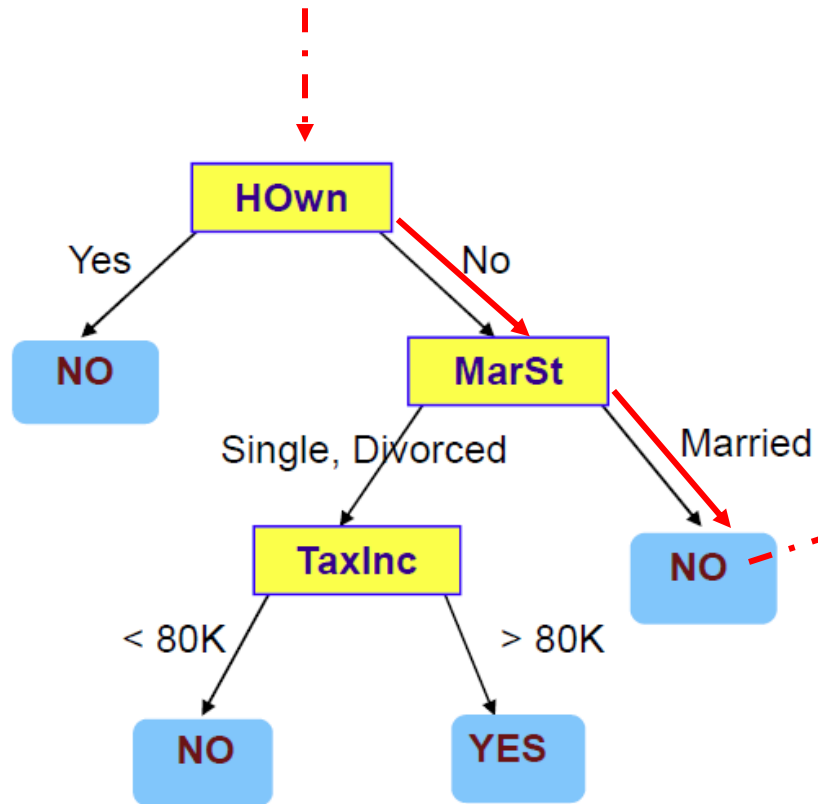
Start from the root of tree



Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Apply Model to Test Data

Start from the root of tree



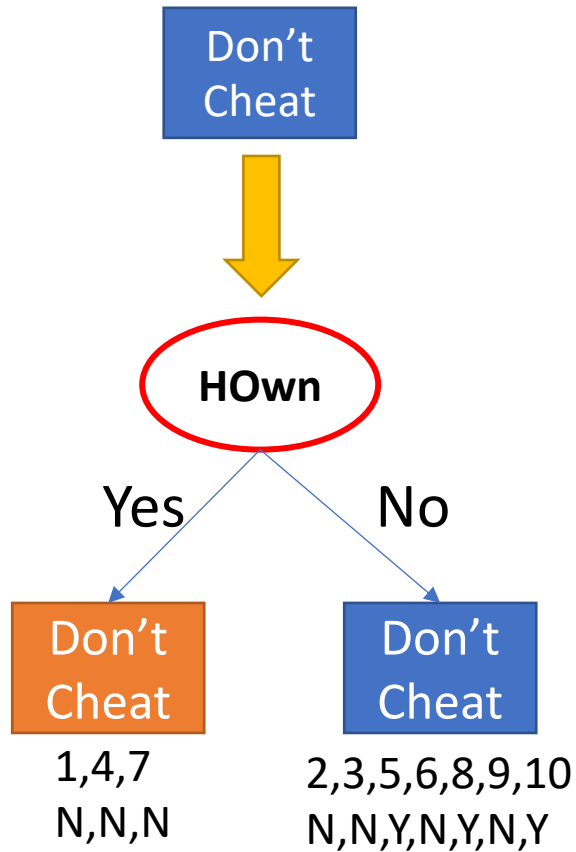
Home Owner	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Assign Cheat to "No"

Hunt's Algorithm

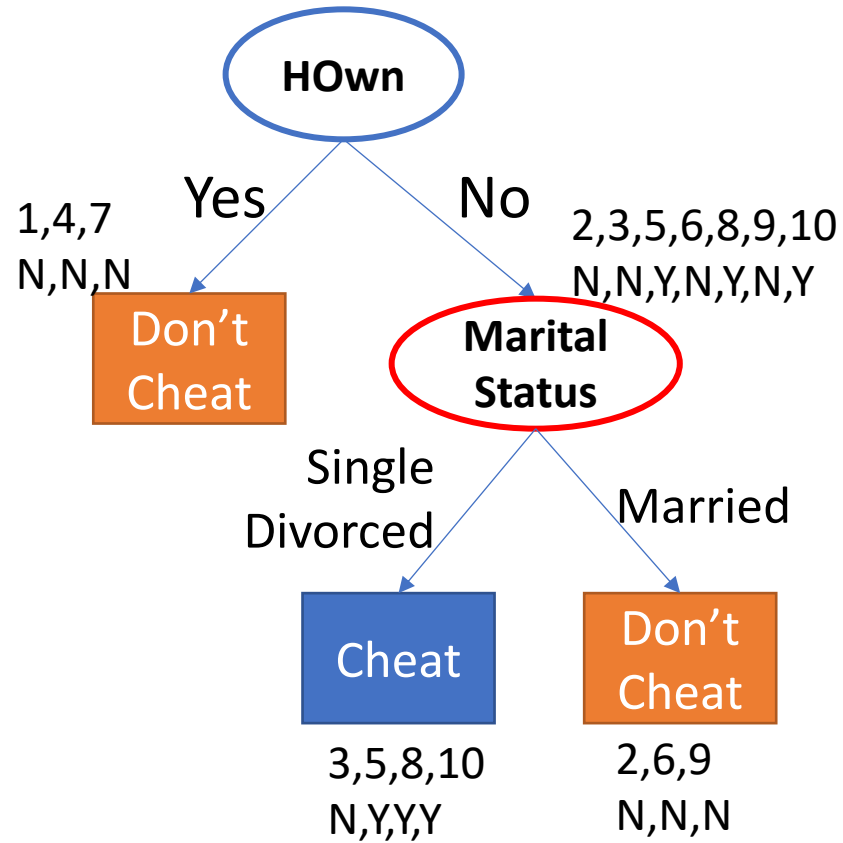
- Let D_t be the set of training records that reach a node t
- General Procedure:
 - If D_t contains records that belong to **the same class** y_t , then t is a leaf node labeled as y_t .
 - If D_t is an **empty set**, then t is a leaf node labeled by the default class y_d
 - If D_t contains records that belong to **more than one class**, use an attribute to split the data into smaller subsets.
 - **Recursively** apply the procedure to each subset.

Hunt's Algorithm



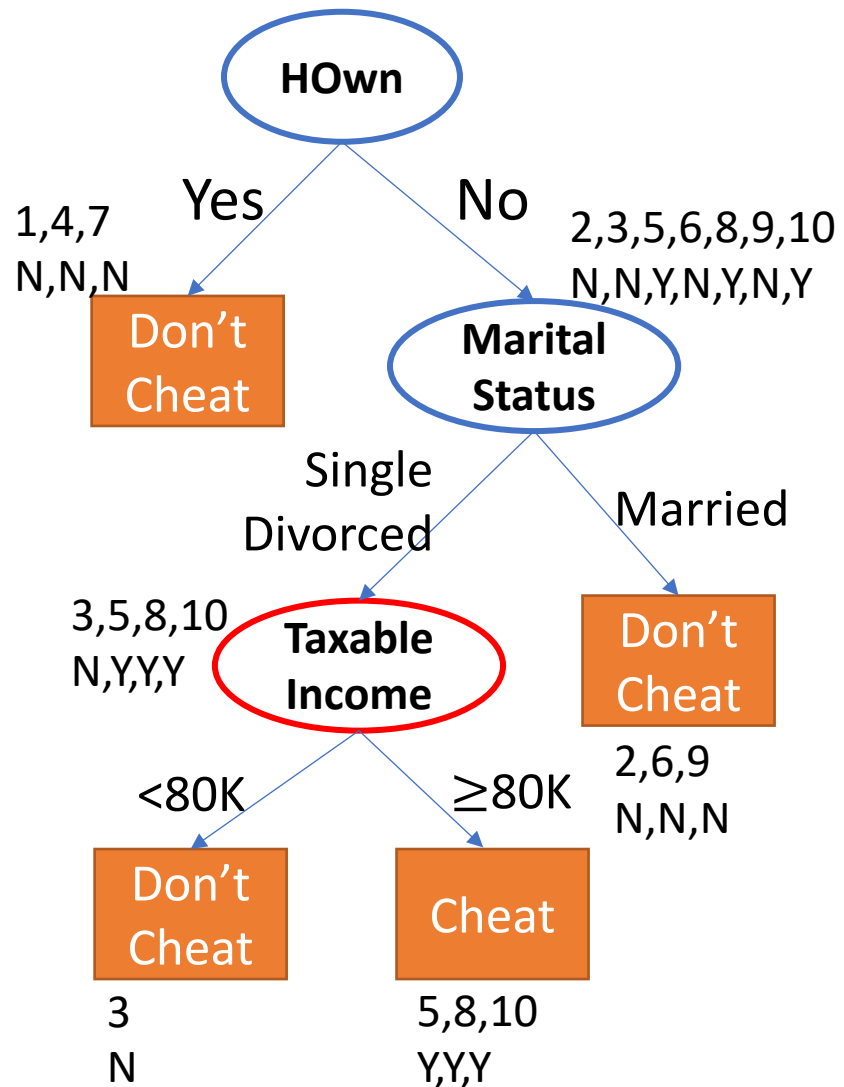
ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Hunt's Algorithm



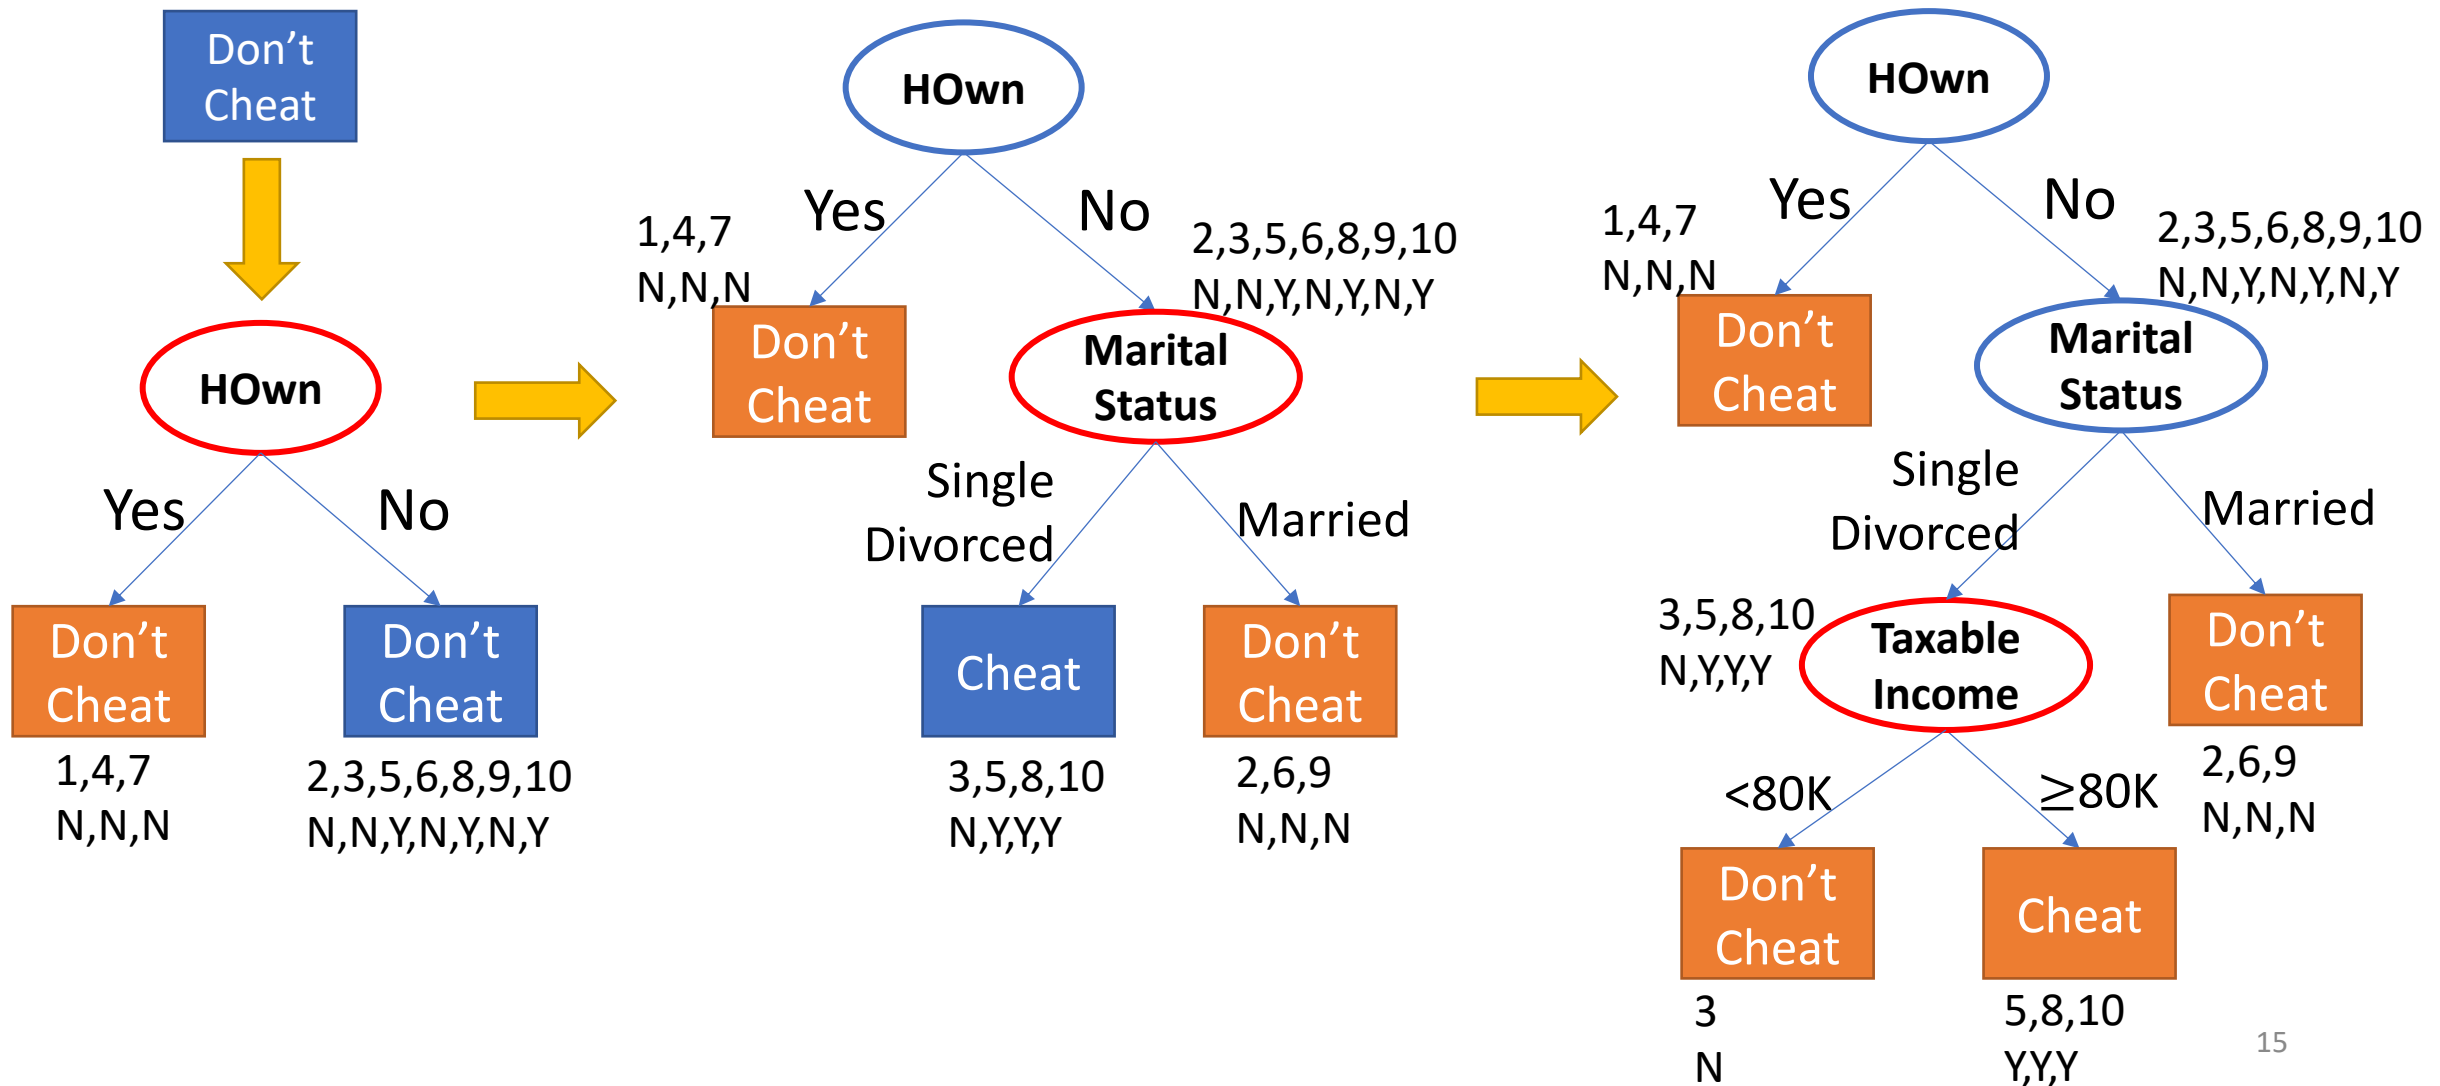
ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Hunt's Algorithm



ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Hunt's Algorithm



How to Split


- Greedy strategy
 - Split based on the training data
 - Split the records based on an attribute that optimizes certain criterion.
- Issues
 - Determine how to split the records
 - How to specify the attribute condition?
 - How to determine the best split?
 - Determine when to stop splitting

How to Specify the Attribute Condition?


- Depends on **attribute types**
 - Nominal
 - Ordinal
 - Continuous
- Depends on **number of ways to split**
 - 2-way split
 - Multi-way split

Splitting Based on Nominal Attributes

- **Multi-way split:** Use as many partitions as distinct values.




- **Binary split:** Divides values into two subsets. Need to find optimal partitioning.




Splitting Based on Ordinal Attributes


- Multi-way split: Use as many partitions as distinct values.



- Binary split: Divides values into two subsets. Need to find optimal partitioning.



- What about this split? Preserve **order** property among attribute values




Splitting Based on Continuous Attributes


Different ways of handling:

- **Discretize** to form an ordinal categorical attribute
 - Ranges can be found by equal interval bucketing, equal frequency bucketing (percentiles), or clustering.
- **Binary Decision**: $(A < v)$ or $(A \geq v)$
 - Consider all possible splits and find the best cut
 - Can be more computing intensive

Splitting Based on Continuous Attributes



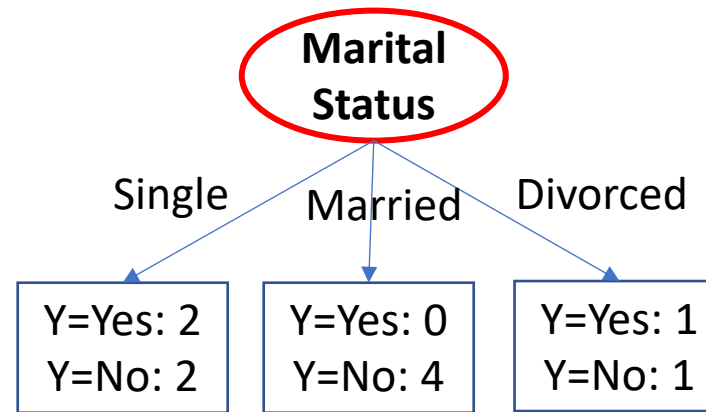

Binary split



Multi-way split

How to Determine the Best Split

- Which attribute we prefer to split on?
- Idea: use counts as leaves to define probability distribution, so we can measure uncertainty.



ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

How to Determine the Best Split

- Greedy approach:
 - Nodes with **homogeneous** class distribution are preferred
- Need a measure of node impurity:

Y=Yes: 5
Y=No: 5

Non-homogeneous
High degree of impurity

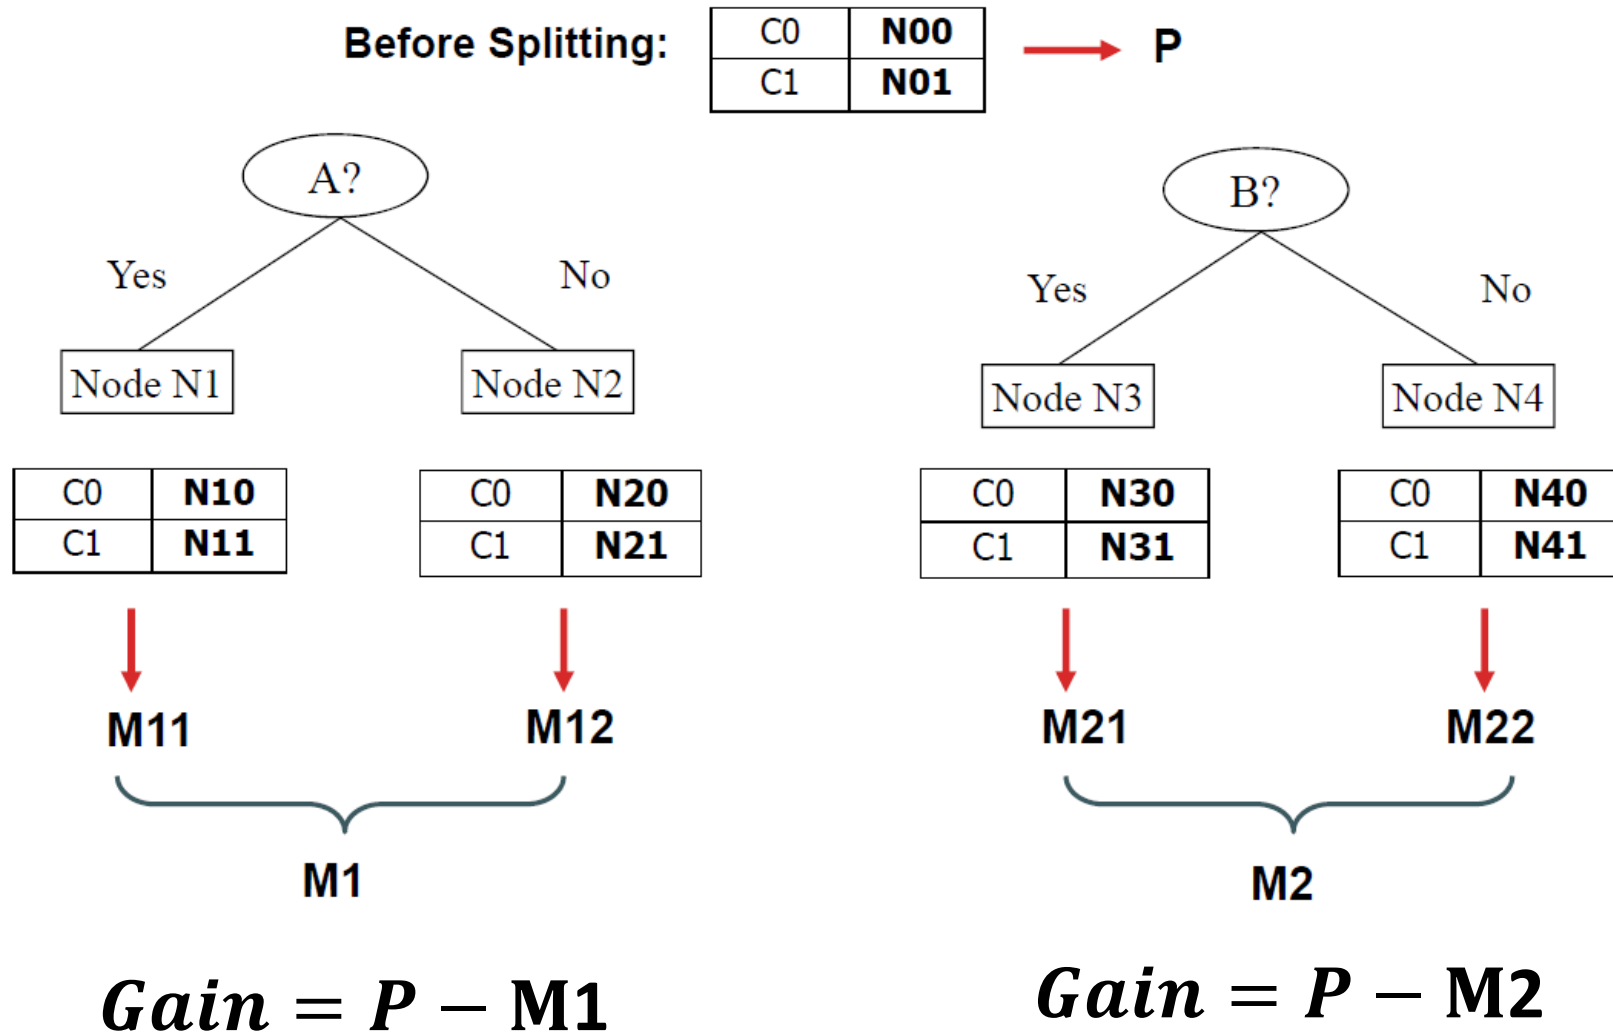
Y=Yes: 9
Y=No: 1

Homogeneous
Low degree of impurity

Measures of Node Impurity

- Gini Index
- Entropy

How to Find the Best Split



How to Find the Best Split

1. Compute impurity measure (P) before splitting
2. Compute impurity measure (M) after splitting
 - Compute impurity measure of each child node
 - Compute the average impurity of the children (M)
3. Choose the attribute test condition that produces the **highest gain**

$$Gain = P - M$$

or equivalently, **lowest impurity measure** after splitting (M)

Measure of Impurity: GINI

- Gini Index for a given node t :
 - $GINI(t) = 1 - \sum_j [p(j|t)]^2$
 - Note: $p(j|t)$ is the relative frequency of class j at node t .
- Maximum $\left(1 - \frac{1}{n_c}\right)$ when records are equally distributed among all classes, implying least interesting information. n_c is the number of classes.
- Minimum (0.0) when all records belong to one class, implying most interesting information.

C1	0
C2	6
Gini=0.000	

C1	1
C2	5
Gini=0.278	

C1	2
C2	4
Gini=0.444	

C1	3
C2	3
Gini=0.500	

Examples for Computing GINI

$$GINI(t) = 1 - \sum_j [p(j|t)]^2$$

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Gini = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444$$

Note: $p(j|t)$ is the relative frequency of class j at node t .

Split Based on GINI

- When a node p is split into k partitions (children), the quality of split is computed as

$$GINI_{split} = \sum_{i=1}^k \frac{n_i}{n} GINI(i)$$


- n_i = number of records at child i ,
- n = number of records at node p .
- Choose the attribute that **minimizes weighted average** Gini index of the children.

Binary Attributes: Computing GINI Index

- Splits into two partitions
- Effect of Weighing partitions: Larger and Purer Partitions needed

	Parent
C1	7
C2	5
Gini = 0.486	

$$\begin{aligned}
 Gini(Parent) &= 1 - \left(\frac{7}{12}\right)^2 - \left(\frac{5}{12}\right)^2 \\
 &= 0.4867
 \end{aligned}$$



	N1	N2
C1	5	2
C2	1	4
Gini=0.361		

$$\begin{aligned}
 Gini(N1) &= 1 - \left(\frac{5}{6}\right)^2 - \left(\frac{1}{6}\right)^2 \\
 &= 0.278
 \end{aligned}$$

$$\begin{aligned}
 Gini(N2) &= 1 - \left(\frac{2}{6}\right)^2 - \left(\frac{4}{6}\right)^2 \\
 &= 0.444
 \end{aligned}$$

$$\begin{aligned}
 Gini(Children) &= \frac{6}{12} \times 0.278 + \frac{6}{12} \times 0.444 \\
 &= 0.361
 \end{aligned}$$

Categorical Attributes: Computing Gini Index

- For each **distinct value**, gather counts for each class in the dataset
- Use the count matrix to make decisions

Multi-way split

	CarType		
	Family	Sports	Luxury
C1	1	8	1
C2	3	0	7
Gini	0.163		

Gini(Multi-way)

$$\begin{aligned} &= (4/20) * (1 - (1/4)^2 - (3/4)^2) + \\ &\quad (8/20) * (1 - (0/8)^2 - (8/8)^2) + \\ &\quad (8/20) * (1 - (1/8)^2 - (7/8)^2) \\ &= 0.163 \end{aligned}$$

Two-way split
(find **best** partition of values)

	CarType	
	{Sports, Luxury}	{Family}
C1	9	1
C2	7	3
Gini	0.468	

	CarType	
	{Sports}	{Family, Luxury}
C1	8	2
C2	0	10
Gini	0.167	


Gini(Two-way)

$$\begin{aligned} &= (16/20) * (1 - (9/16)^2 - (7/16)^2) + \\ &\quad (4/20) * (1 - (1/4)^2 - (3/4)^2) \\ &= 0.468 \end{aligned}$$

Continuous Attributes: Computing Gini Index

- Use **Binary Decisions** based on **one value**
- Several Choices for the splitting value
 - Possible splitting values: distinct values
- Each splitting value has a **count matrix**
 - Class counts in each of the partitions,
 $A < v$ and $A \geq v$
- Simple method to choose best v
 - For each v , scan the database to gather count matrix and compute its Gini index
 - Computationally inefficient! Repetition of work.

ID	Home Owner	Marital Status	Annual Income	Defaulted Borrower
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



Entropy

- Entropy for a given node t :
 - $\text{Entropy}(t) = -\sum_j p(j|t) \log_2 p(j|t)$
 - Note: $p(j|t)$ is the relative frequency of class j at node t .
- Measures homogeneity of a node.
 - Maximum ($\log_2 n_c$) when records are equally distributed among all classes, implying least interesting information.
 - Minimum (0.0) when all records belong to one class, implying most interesting information.
- Entropy based computations are similar to the GINI index computations

Examples for Computing Entropy

$$\text{Entropy}(t) = - \sum_j p(j|t) \log_2 p(j|t)$$

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$\text{Entropy} = -0 \log 0 - 1 \log 1 = -0 - 0 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$\text{Entropy} = - (1/6) \log_2 (1/6) - (5/6) \log_2 (5/6) = 0.65$$

C1	3
C2	3

$$P(C1) = 3/6 \quad P(C2) = 3/6$$

$$\text{Entropy} = - (3/6) \log_2 (3/6) - (3/6) \log_2 (3/6) = 1$$




Note: $p(j|t)$ is the relative frequency of class j at node t .

Splitting Based on Information Gain

- Information Gain:
 - $Gain_{split} = Entropy(p) - \sum_{i=1}^k \frac{n_i}{n} Entropy(i)$
 - Parent Node p is split into k partitions;
 - n_i is number of records in partition i
- Measures reduction in Entropy achieved because of the split.
- Choose the split that achieves **most reduction** (maximizes GAIN).
- Disadvantage:
 - Tends to prefer splits that result in **large number of partitions**, each being small but pure.

Problems with Information Gain

- Information gain tends to prefer splits that result in large number of partitions, each being small but pure.
- Customer ID** has the highest information gain because entropy for all the children is zero.



Splitting Based on Gain Ratio

- Gain Ratio:

- $\text{GainRatio}_{split} = \frac{\text{Gain}_{split}}{\text{SplitINFO}}, \quad \text{SplitINFO} = -\sum_{i=1}^k \frac{n_i}{n} \log \frac{n_i}{n}$

- Parent Node p is split into k partitions;
- n_i is number of records in partition i
- Adjusts information gain by the entropy of the partitioning (SplitINFO). Higher entropy partitioning (large number of small partitions) is penalized!
- Designed to overcome the disadvantage of Information Gain
 - For example : 10 records (parent node)
 - Two equal partitions (5,5): $\text{SplitINFO} = -2 \cdot (5/10) \log_2(5/10) = 1$
 - Ten equal partitions: $\text{SplitINFO} = -10 \cdot (1/10) \log_2(1/10) = 3.32$

Decision Tree Based Classification

- Advantages:
 - Not computationally expensive to construct
 - Extremely fast at classifying unknown records
 - Easy to interpret for small-sized trees
 - Accuracy is comparable to other classification techniques for many simple datasets
- Disadvantages: **Overfitting**
 - Overfitting due to lack of representative samples or due to some noise
 - Overfitting results in decision trees that are more complex than necessary
 - Training error no longer provides a good estimate of how well the tree will perform on previously unseen records

How to Address Overfitting: Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
 - Stop if all instances belong to the same class
 - Stop if all the attribute values are the same
- More restrictive conditions:
 - Stop if number of instances is less than some user-specified threshold
 - Stop if class distribution of instances are independent of the available features
 - Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

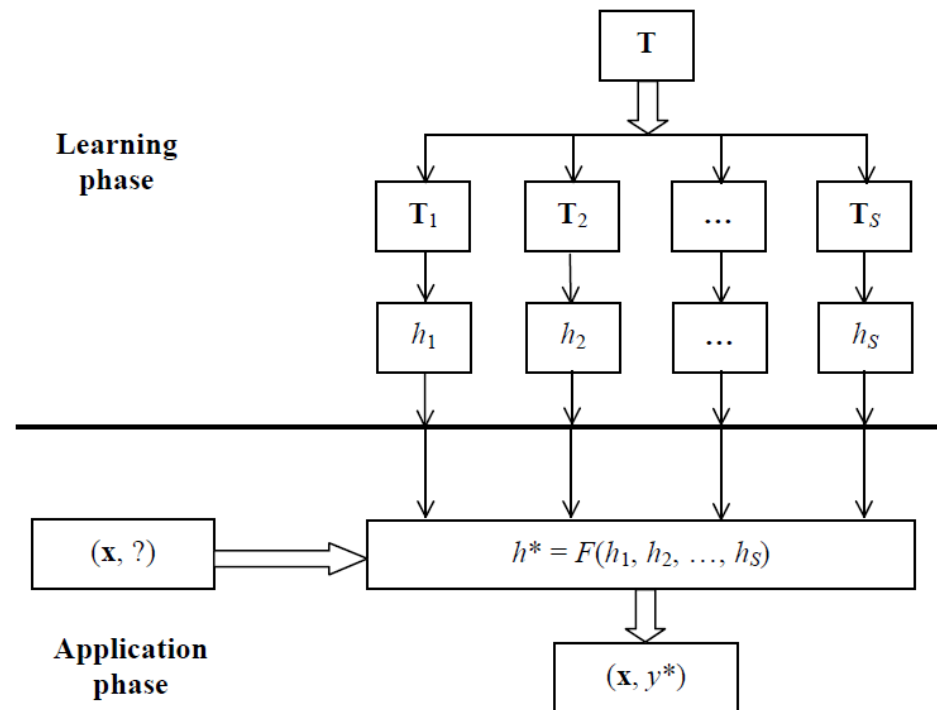
How to Address Overfitting: Post-Pruning

- Grow decision tree to its entirety
- Trim the nodes of the decision tree in a bottom-up fashion
- If generalization error improves after trimming, replace sub-tree by a leaf node.
- Class label of leaf node is determined from majority class of instances in the sub-tree.

Ensemble Methods

Ensemble

- Basic idea: Combine multiple models into one!
 - Build different “experts” and let them vote




Different training sets
and/or learning algorithms

Motivations of Ensemble Methods

- Improve the accuracy and robustness over single model methods
- Efficiency: a complex problem can be decomposed into multiple sub-problems that are easier to understand and solve (divide-and-conquer approach)
- Applications:
 - Distributed computing
 - Privacy-preserving applications
 - Large-scale data with reusable models
 - Multiple sources of data

Why Ensemble Works?

- The given task may be too complex, or lie outside the space of functions that can be implemented by the chosen classifier method
- Appropriate combinations of simple (e.g., linear) classifiers can learn complex (e.g., non-linear) boundaries
- **Ensemble gives the global picture!**



Pros and Cons

- Advantages:
 - Improve predictive performance
 - Different types of classifiers can be directly included
 - Easy to implement
 - Not too much parameter tuning
- Disadvantages:
 - The combined classifier is not transparent (black box)
 - Not a compact representation

How to Make an Effective Ensemble?

- Two basic decisions when designing ensembles:
 - How to generate the base classifiers?
 - How to integrate/combine them?

Generating Base Classifiers

- Sampling training examples
 - Train k classifiers on k subsets drawn from the training set
- Using different learning models
 - Use all the training examples, but apply different learning algorithms
- Sampling features
 - Train k classifiers on k subsets of features drawn from the feature space
- Learning “randomly”
 - Introduce randomness into learning procedures

How to Integrate?

- Majority voting
- Weighted majority voting

Diversity and Accuracy

- The individual classifiers must be diverse (errors on different data)
- If they make the same errors, such mistakes will be carried into the final prediction
- The base classifiers need to be “reasonably accurate” to avoid poor classifiers to obtain the majority of votes.

Ensemble Methods

- Predict class label for unseen data by aggregating a set of predictions (classifiers learned from the training data)
 - **Bagging** (Breiman 1994 “Bagging Predictors”)
 - **Boosting** (Freund and Schapire 1995, Friedman et al. 1998)
 - **Random forests** (Breiman 2001 “Random Forests”)

Bagging: Bootstrap Aggregation


- Take repeated **bootstrap** samples from training set D (Breiman, 1994)
- **Bootstrap sampling:**
 - Bootstrap sampling: Given set D containing N training examples, create D' by drawing N examples at random with replacement from D
- Bagging:
 - Create k bootstrap samples D_1, \dots, D_k
 - Train distinct classifier on each D_i
 - Classify new instances by majority vote/average

Bagging Example

- Only one feature

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
Y	1	1	1	-1	-1	-1	-1	1	1	1

- Decision stump



Bagging Example: Classifier 1-5

Bagging Round 1:

x	0.1	0.2	0.2	0.3	0.4	0.4	0.5	0.6	0.9	0.9
y	1	1	1	1	-1	-1	-1	-1	1	1

$x \leq 0.35 \implies y = 1$

$x > 0.35 \implies y = -1$

Bagging Round 2:

x	0.1	0.2	0.3	0.4	0.5	0.8	0.9	1	1	1
y	1	1	1	-1	-1	1	1	1	1	1

$x \leq 0.65 \implies y = 1$

$x > 0.65 \implies y = 1$

Bagging Round 3:

x	0.1	0.2	0.3	0.4	0.4	0.5	0.7	0.7	0.8	0.9
y	1	1	1	-1	-1	-1	-1	-1	1	1

$x \leq 0.35 \implies y = 1$

$x > 0.35 \implies y = -1$

Bagging Round 4:

x	0.1	0.1	0.2	0.4	0.4	0.5	0.5	0.7	0.8	0.9
y	1	1	1	-1	-1	-1	-1	-1	1	1

$x \leq 0.3 \implies y = 1$

$x > 0.3 \implies y = -1$

Bagging Round 5:

x	0.1	0.1	0.2	0.5	0.6	0.6	0.6	1	1	1
y	1	1	1	-1	-1	-1	-1	1	1	1

$x \leq 0.35 \implies y = 1$

$x > 0.35 \implies y = -1$

Bagging Example: Classifier 6-10

Bagging Round 6:

x	0.2	0.4	0.5	0.6	0.7	0.7	0.7	0.8	0.9	1
y	1	-1	-1	-1	-1	-1	-1	1	1	1

$x \leq 0.75 \implies y = -1$

$x > 0.75 \implies y = 1$

Bagging Round 7:

x	0.1	0.4	0.4	0.6	0.7	0.8	0.9	0.9	0.9	1
y	1	-1	-1	-1	-1	1	1	1	1	1

$x \leq 0.75 \implies y = -1$

$x > 0.75 \implies y = 1$

Bagging Round 8:

x	0.1	0.2	0.5	0.5	0.5	0.7	0.7	0.8	0.9	1
y	1	1	-1	-1	-1	-1	-1	1	1	1

$x \leq 0.75 \implies y = -1$

$x > 0.75 \implies y = 1$

Bagging Round 9:

x	0.1	0.3	0.4	0.4	0.6	0.7	0.7	0.8	1	1
y	1	1	-1	-1	-1	-1	-1	1	1	1

$x \leq 0.75 \implies y = -1$

$x > 0.75 \implies y = 1$

Bagging Round 10:

x	0.1	0.1	0.1	0.1	0.3	0.3	0.8	0.8	0.9	0.9
y	1	1	1	1	1	1	1	1	1	1

$x \leq 0.05 \implies y = -1$

$x > 0.05 \implies y = 1$

Bagging Example

Round	x=0.1	x=0.2	x=0.3	x=0.4	x=0.5	x=0.6	x=0.7	x=0.8	x=0.9	x=1.0
1	1	1	1	-1	-1	-1	-1	-1	-1	-1
2	1	1	1	1	1	1	1	1	1	1
3	1	1	1	-1	-1	-1	-1	-1	-1	-1
4	1	1	1	-1	-1	-1	-1	-1	-1	-1
5	1	1	1	-1	-1	-1	-1	-1	-1	-1
6	-1	-1	-1	-1	-1	-1	-1	1	1	1
7	-1	-1	-1	-1	-1	-1	-1	1	1	1
8	-1	-1	-1	-1	-1	-1	-1	1	1	1
9	-1	-1	-1	-1	-1	-1	-1	1	1	1
10	1	1	1	1	1	1	1	1	1	1
Sum	2	2	2	-6	-6	-6	-6	2	2	2
Sign	1	1	1	-1	-1	-1	-1	1	1	1
True Class	1	1	1	-1	-1	-1	-1	1	1	1

Figure 5.36. Example of combining classifiers constructed using the bagging approach.

Boosting

Basic idea:

- Assign a weight to every training set instance
- **Initially**, all instances have the **same weight**
- As the boosting proceeds, adjust weights based on how well we have predicted data points so far
 - data points **correctly** predicted → **low** weight
 - data points **mis-predicted** → **high** weight
- Results: as learning proceeds, the learner is forced to focus on portions of data space not previously well predicted

Formal Description of Boosting

- Given training set $(x_1, y_1), \dots, (x_N, y_N)$
- $y_i \in \{-1, +1\}$ correct labels of instance $x_i \in X$
- For $t = 1, \dots, T$:
 - Construct weight distribution $D^t(i)$ on $\{1, \dots, N\}$
 - Find weak classifier $h_t: X \rightarrow \{-1, +1\}$
 - With error ϵ_t on $D^t(i)$: $\epsilon_t = P_{i \sim D^t} [h_t(x_i) \neq y_i]$
- Output final/combined classifier H_{final}

AdaBoost

- Given: $(x_1, y_1), \dots, (x_N, y_N)$ where $x_i \in X, y_i \in \{-1, +1\}$
- Initialize $D_1(i) = \frac{1}{N}$
- For $t = 1, \dots, T$:
 - Train weak learner using distribution D_t Naïve Bayes, decision stump,...
 - Get weak classifier $h_t: X \rightarrow \{-1, +1\}$
 - Choose weight for the classifier: $\alpha_t \in \mathbb{R}$. Increase weight if predicting incorrectly
 - Update: $D_t(i) = \frac{D_t(i)}{Z_t} \begin{cases} e^{-\alpha_t} & \text{if } y_i = h_t(x_i) \\ e^{\alpha_t} & \text{if } y_i \neq h_t(x_i) \end{cases} = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$
with Z_t as a normalization factor $Z_t = \sum_{i=1}^N D_t(i) \exp(-\alpha_t y_i h_t(x_i))$.
- Output the final classifier: $H_{final}(x) = \text{sign}(\sum_{t=1}^T \alpha_t h_t(x))$

How to Choose α_t for Classifier h_t


$$\alpha_t = \frac{1}{2} \ln \left(\frac{1 - \epsilon_t}{\epsilon_t} \right)$$

$$\epsilon_t = P_{i \sim D^{(t)}}[h_t(x_i) \neq y_i] = \sum_{i=1}^N D_t(i) \delta(h_t(x_i) \neq y_i)$$

- $\epsilon_t = 0, \alpha_t = \infty$: if h_t perfectly classifies all weighted data points
- $\epsilon_t = 1, \alpha_t = -\infty$: if h_t classifies incorrectly on all points
- $\epsilon_t = 0.5, \alpha_t = 0$

Smaller error rate, larger weight for voting!


Example of Adaboost



$$D_1(1) = \dots = D_1(10) = \frac{1}{10}$$


weak classifiers = vertical or horizontal half-planes

Example of Adaboost: Round 1




$$err_1 = 0.1 \times 3 = 0.3; \alpha_1 = \frac{1}{2} \times \log\left(\frac{1-0.3}{0.3}\right) = 0.42$$

$$D_t(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$$

For those classified not correctly:	$D_2(i) = 0.1 \times e^{0.42 \times 1} = 0.1527$	Normalization 	0.1667
For those classified correctly:	$D_2(i) = 0.1 \times e^{0.42 \times (-1)} = 0.0654$		0.0714


Example of Adaboost: Round 2



$$err_2 = 0.0714 \times 3 = 0.21; \alpha_2 = \frac{1}{2} \times \log\left(\frac{1-0.21}{0.21}\right) = 0.6625$$

For those classified correctly:	$D_3(i) = 0.1667 \times e^{0.6625 \times (-1)} = 0.0859$		0.1047
For those classified not correctly	$D_3(i) = 0.0714 \times e^{0.6625 \times 1} = 0.1385$	Normalization →	0.1688
For those classified correctly:	$D_3(i) = 0.0714 \times e^{0.6625 \times (-1)} = 0.0368$		0.0449

Example of Adaboost: Round 3



$$err_3 = 0.0449 \times 3 = 0.1347; \alpha_3 = \frac{1}{2} \times \log\left(\frac{1 - 0.1347}{0.1347}\right) = 0.92$$

Final Classifier

$$H_{\text{final}} = \text{sign} \left(0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \right)$$

=

Voted Combination of Classifiers

- The general problem here is to try to combine many simple “weak” classifiers into a single “strong” classifier.

- We consider voted combinations of simple binary component classifiers

$$H_{final}(x) = \alpha_1 h(x; \theta_1) + \cdots + \alpha_T h(x; \theta_T)$$

- Where θ is the model parameter and the non-negative votes α_i can be used to emphasize base classifiers that are more reliable than others.

Random Forests

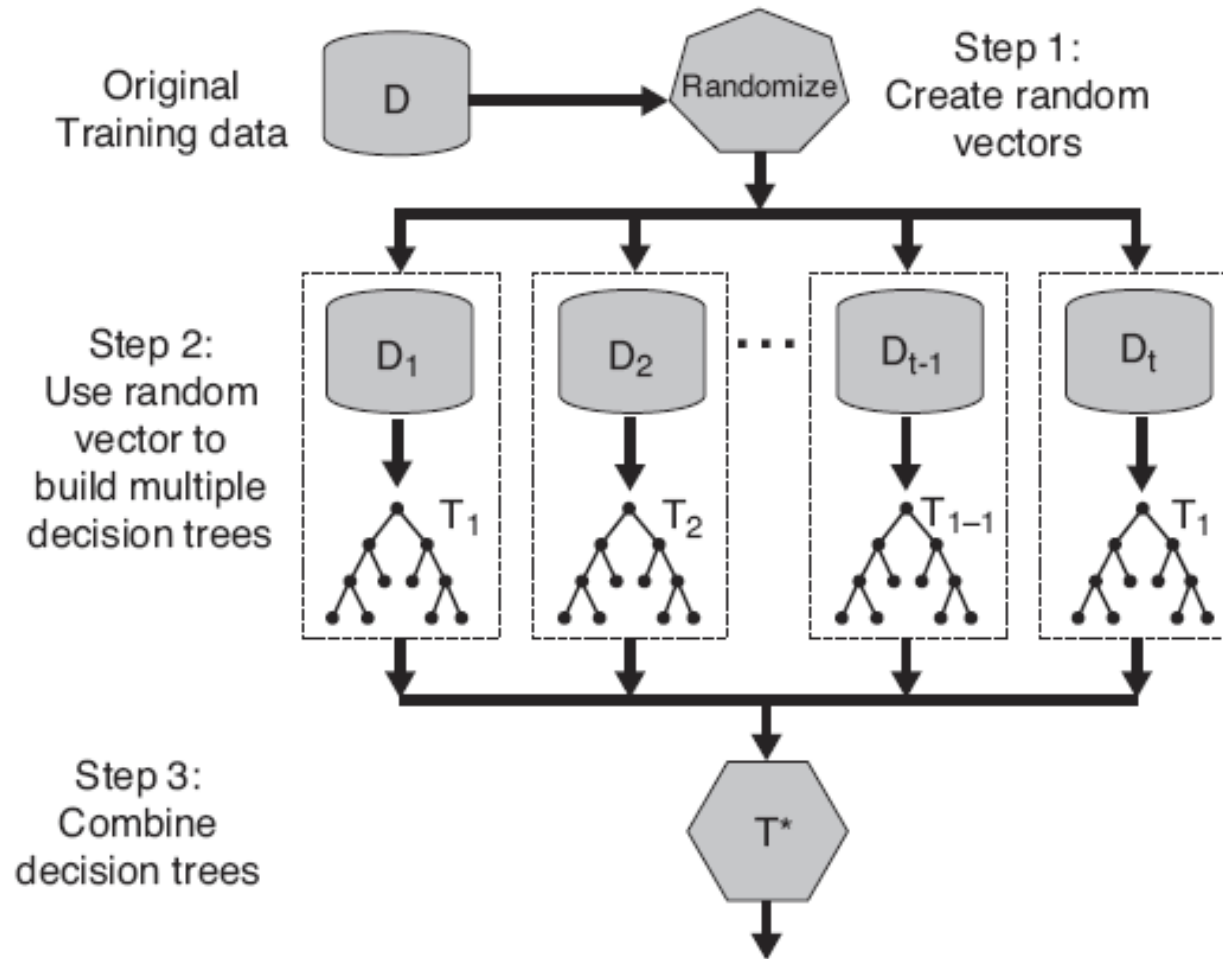
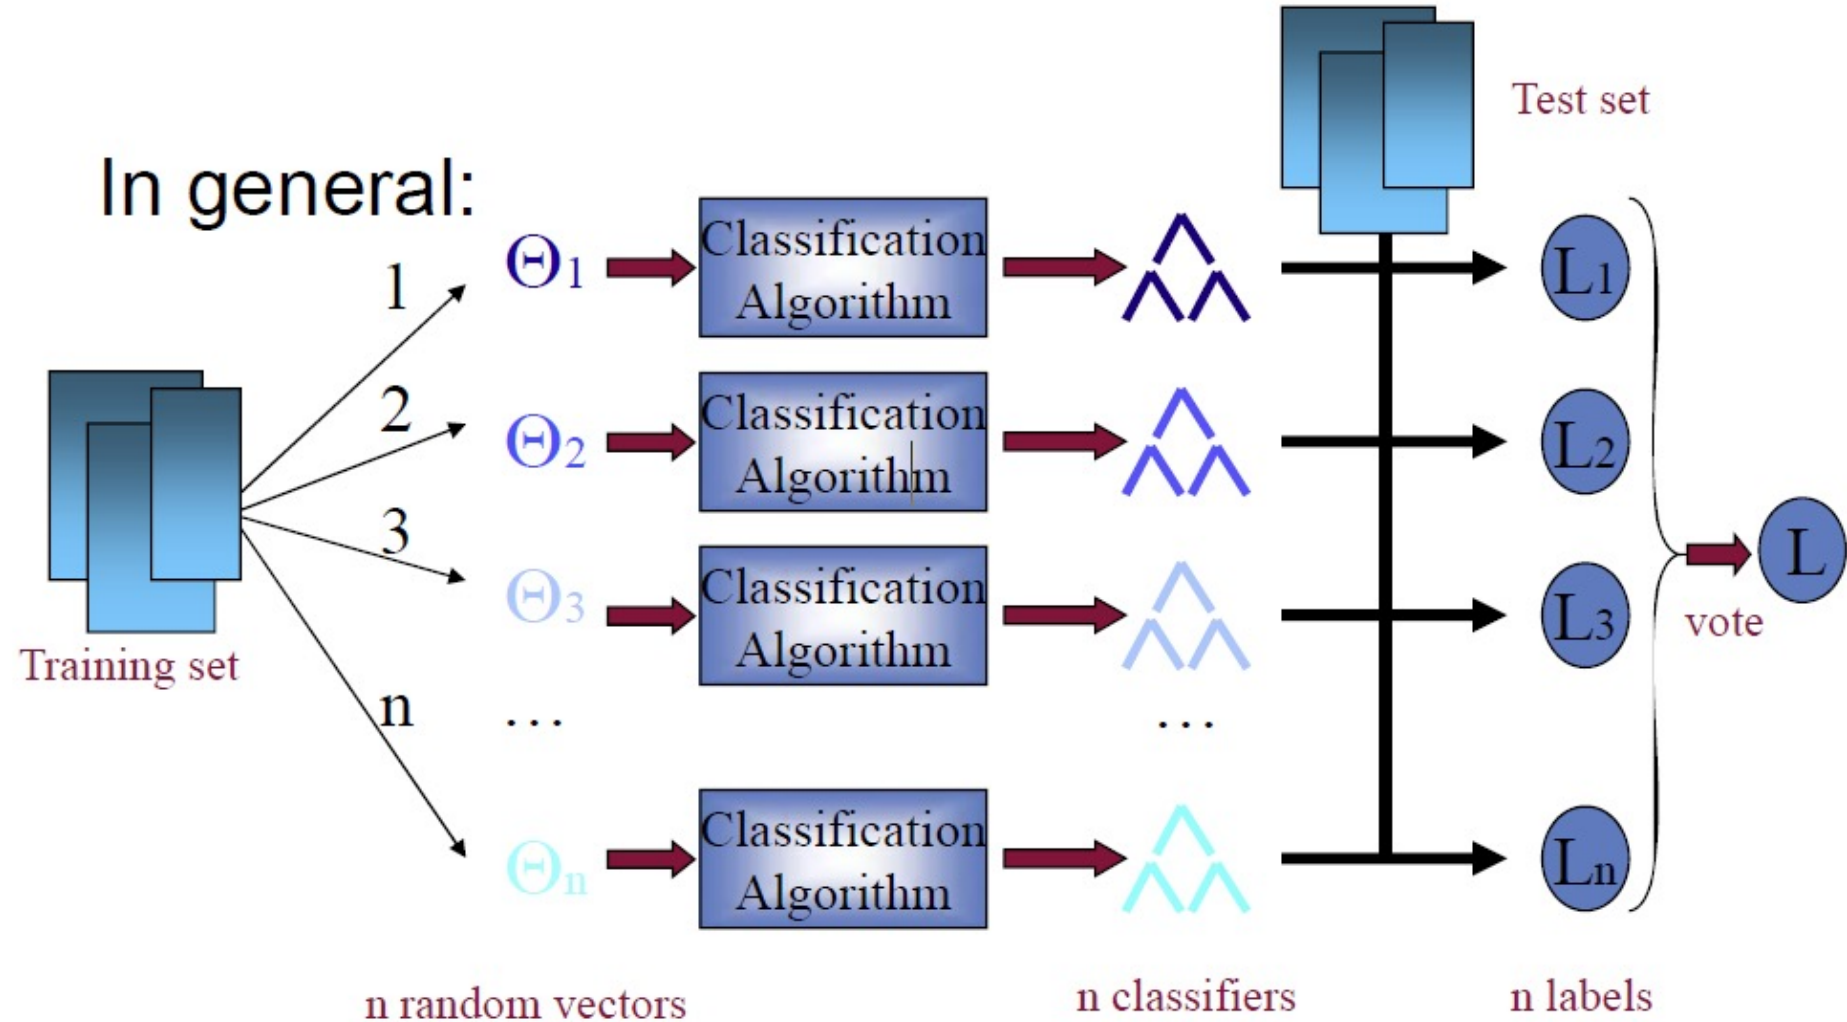


Figure 5.40. Random forests.

Random Forests



Random Forests

- Ensemble method specifically designed for decision tree classifiers
- Two sources of randomness: “bagging” and “random input vectors”
- Use bootstrap aggregation to train many decision trees
 - Randomly subsample N examples
 - Train decision tree on subsample
 - Use average or majority vote among learned trees as prediction
- Also randomly subsample features: best split at each node is chosen from a random sample of m attributes instead of all attributes

Random Forests Algorithm

- For $b = 1$ to B
 - Draw a bootstrap sample of size N from the data
 - Grow a tree T_b using the bootstrap sample as follows
 - Choose m attributes uniformly at random from the data
 - Choose the best attribute among the m to split on
 - Split on the best attribute and recurse until partitions have fewer than s_{min} number of nodes
- Prediction for a new data point x
 - **Regression**: $\frac{1}{B} \sum_b T_b(x)$
 - **Classification**: choose the majority class label among $T_1(x), \dots, T_B(x)$.

Readings

1. <https://hunch.net/~coms-4771/quinlan.pdf>
2. <https://arxiv.org/abs/1603.02754>
3. <https://www.jstor.org/stable/2699986?refreqid=excelsior%3A4a12047a831ef2f51e03a13d2a4e52ee>

Summary of Today's Lecture

- Decision Tree
- Ensemble Methods
 - Bagging
 - Boosting
 - Random Forests