

# CS 559 Machine Learning

## Lecture 9: Clustering, K-means, and GMM

Ping Wang

Department of Computer Science

Stevens Institute of Technology




# Today's Lecture

- Clustering
- K-means Algorithm
- Gaussian Mixture Model (GMM)

# What is Clustering?

- Finding groups of objects such that the objects in a group will be similar (or related) to one another and different from (or unrelated to) the objects in other groups.



# Applications of Clustering

- Group related documents for browsing
- Group genes and proteins that have similar functionality
- Group stocks with similar price fluctuations
- Group the regions with similar temperature
- And many others...

# What is Not Clustering?

- **Simple segmentation**

- Dividing students into different registration groups alphabetically, by last name.


- **Results of a query**

- Groupings are a result of an external specific query conditions.
  - Clustering is a grouping of objects based on the data.



- **Supervised classification**

- Have class label information



# Notion of a Cluster can be Ambiguous





How many clusters?



Six Clusters



Two Clusters



Four Clusters



# Distance Based Clustering

- Fundamental to all clustering techniques is the measure of **distance or dissimilarity** between two objects.
- **What is Similarity?:**
  - The quality or state of being similar;
  - Likeness; resemblance; a similarity of features - Websters' dictionary.
  - Similarity is hard to define, but we know it when we see it.
  - Here, we use the distance to measure the similarity/dissimilarity.

# Dissimilarities Based on Features

- Assuming  $x^i = (x_1^i, x_2^i, \dots, x_d^i)^T \in \mathbb{R}^d, i = 1, \dots, N$
- Distance (dissimilarity):  $D(x^i, x^j) = \sum_{k=1}^d d_k(x_k^i, x_k^j)$
- Squared **Euclidean distance**:

$$d_k(x_k^i, x_k^j) = (x_k^i - x_k^j)^2$$

$$\Rightarrow D(x^i, x^j) = \sum_{k=1}^d (x_k^i - x_k^j)^2$$



# Non-Probabilistic Algorithms

- These algorithms work **directly on the observed data**, without relying on a probability model about the data.
- Commonly used in data mining, since usually the prior knowledge about the data generation is not available.

# Non-Probabilistic Algorithms

- Assuming  $x_i \in \mathbb{R}^d, i = 1, \dots, N$
- Pre-specified number of cluster  $K, k \in \{1, \dots, K\}$
- Each data point  $x_i$  is assigned to **one, and only one cluster**.
- **Goal:** Divide the data into  $K$  clusters to satisfy a required objective, defined by a **dissimilarity function  $D(x^i, x^j)$** .
- Usually, the assignment of data to clusters is done by **minimizing a “loss” function** that measures the degrees to which the clustering goal is not met.

# Non-probabilistic Algorithms

- Since the goal is to assign close points to the same cluster, a natural loss function would be **within cluster scatter**. Need to define the following two notations.
- Notations:
  - **Binary indicator  $r_{nk}$** : (describe which of the  $K$  clusters the data point  $x_n$  is assigned to)

$$r_{nk} = \begin{cases} 1, & \text{if } x_n \text{ is assigned to cluster } k \\ 0, & \text{otherwise} \end{cases}$$

E.g., if  $K = 3$ ,  $x_n$  is assigned to cluster 2, then  $r_{n1} = 0, r_{n2} = 1, r_{n3} = 0$ . Known as 1-of-K coding scheme.

- **Centroid  $\mu_k$  for each cluster**:  $\mu_k \in \mathcal{R}^d$  and can be considered as centers of clusters.

# Objective Function $J$

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} D(x_n, \mu_k)$$

- Also called **distortion measure**.
- Represents the sum of distances of each data point to the center of its assigned cluster  $\mu_k$ .
- Need to select  $D(.,.)$ , find  $r_{nk}$  and  $\mu_k$ .

# Objective of K-means Clustering

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} D(x_n, \mu_k)$$

- K-means is one of **the most popular iterative descent** clustering methods.
- Use squared **Euclidean distance** as  $D(.,.)$ :

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

- Represents the sum of square distance of each data point to its assigned vector  $\mu_k$
- Features: quantitative type.
- Goal is to find values for the  $r_{nk}$  and the  $\mu_k$  that can minimize  $J$ .
- Through iterative procedure...

# K-means Algorithm

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

- **Minimize  $J$**  in an iterative procedure to find out  $r_{nk}$  and  $\mu_k$ :
  - **First phase: cluster assignment step**
    - Choose some initial values for the  $\mu_k$ . Then minimize  $J$  with respect to the  $r_{nk}$ , keeping the  $\mu_k$  fixed.
  - **Second phase: centroid update step**
    - We minimize  $J$  with respect to the  $\mu_k$ , keeping  $r_{nk}$  fixed.

# K-means Algorithm: Cluster Assignment Step

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

- **First phase:** minimize  $J$  with respect to the  $r_{nk}$ , keeping the  $\mu_k$  fixed.
  - When  $\mu_k$  fixed,  $J$  becomes **linear function of  $r_{nk}$** .
  - Terms involving different  $n$  are independent, optimize for each  $n$  separately.
  - Choosing  $r_{nk}$  to be 1 for the value of  $k$  that gives the minimum value of  $\|x_n - \mu_k\|^2$  (i.e., assign  $x_n$  to the closest  $\mu_k$ ).

$$r_{nk} = \begin{cases} 1, & \text{if } k = \arg \min_j \|x_n - \mu_j\|^2 \\ 0, & \text{otherwise} \end{cases}$$

# K-means Algorithm: Centroid Update Step

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

- **Second phase:** We minimize  $J$  with respect to the  $\mu_k$ , keeping  $r_{nk}$  fixed.

- $J$  is quadratic function of  $\mu_k$ .

- Let  $\frac{\partial J}{\partial \mu_k} = 0$ :

$$2 \sum_{n=1}^N r_{nk} (x_n - \mu_k) = 0$$
$$\mu_k = \frac{\sum_n r_{nk} x_n}{\sum_n r_{nk}}$$

- Mean of all of the data points  $x_n$  assigned to cluster  $k$ .



# K-means Algorithm


- Summary: initialize the  $\mu_k$ , then iterate the following **two steps** until convergence.
  - **Cluster Assignment Step**: Minimize  $J$  with respect to the  $r_{nk}$ , keeping the  $\mu_k$  fixed.

$$r_{nk} = \begin{cases} 1, & \text{if } k = \arg \min_j \|x_n - \mu_j\|^2 \\ 0, & \text{otherwise} \end{cases}$$

- **Centroid update step**: Minimize  $J$  with respect to the  $\mu_k$ , keeping  $r_{nk}$  fixed.


$$\mu_k = \frac{\sum_n r_{nk} x_n}{\sum_n r_{nk}}$$

# K-means: Example




Initialization: two randomly selected centroid

# K-means: Example




Step 1: assign clusters, determine  $r_{nk}$ .

# K-means: Example




Step 1: recalculate means of clusters, determine  $\mu_k$ .

# K-means: Example




Step 2: assign clusters to new means, determine  $r_{nk}$ .

# K-means: Example




Step 2: recalculate means of clusters, determine  $\mu_k$ .

# K-means: Example



Step 3: assign clusters to new means, determine  $r_{nk}$ .


# K-means: Example



Step 3: recalculate means of clusters, determine  $\mu_k$ .




# K-means: Example



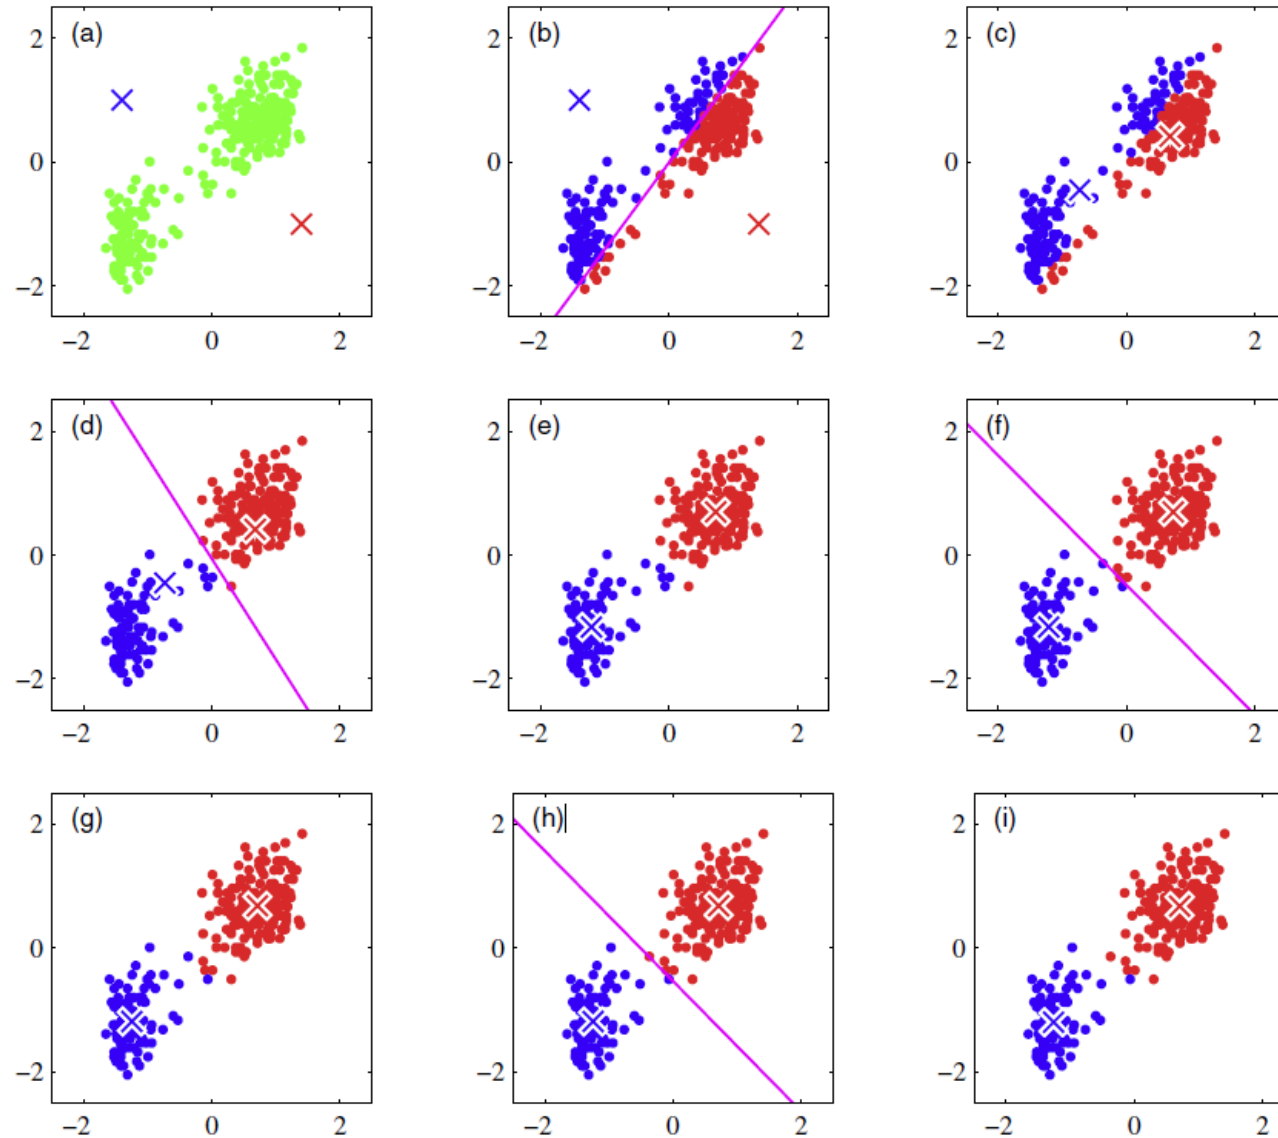
Step 4: assign clusters to new means, determine  $r_{nk}$ .

# K-means: Example



Step 4: recalculate means of clusters, determine  $\mu_k$ . Converged.

# K-means: Example



# K-means Clustering: Pseudocode

---

**Algorithm 16.1:** K-means Algorithm

---

K-means ( $D, k, \epsilon$ ):

```
1  $t = 0$ 
2 Randomly initialize  $k$  centroids:  $\mu_1^t, \mu_2^t, \dots, \mu_k^t$ 
3 repeat
4    $t = t + 1$ 
5   // Cluster Assignment Step
6   foreach  $x_j \in D$  do
7      $j^* = \arg \min_i \{\|x_j - \mu_i^t\|^2\}$  // Assign  $x_j$  to closest centroid
8      $C_{j^*} = C_{j^*} \cup \{x_j\}$ 
9   // Centroid Update Step
10  foreach  $i = 1$  to  $k$  do
11     $\mu_i^t = \frac{1}{|C_i|} \sum_{x_j \in C_i} x_j$ 
12 until  $\sum_{i=1}^k \|\mu_i^t - \mu_i^{t-1}\|^2 \leq \epsilon$ 
```

---

- $D$ : the input data
- $k$ : # of clusters
- $\epsilon$ : error rate
- $t$ : # of iterations


# K-means: Properties and Limitations

- The algorithm converges to **a local minimum**.
- The solution depends on the **initial values**.
- One should start the algorithm with many different random choices for the initial means, and choose the solution having smallest value of the objective function
- The number of clusters  $K$  need to be pre-specified, but can be estimated from the data.
- The algorithm is **sensitive to outliers**.
- Every data point is assigned to one, and only one of the clusters. (**hard assignment**, no probability involved to reflect the level of uncertainty)
- How about probabilistic approach that obtain the **soft assignment** of data points?

# Gaussian Mixture Model and EM

# Motivation of Mixture of Gaussians


- Gaussian distribution has some important analytical properties.
- It also suffers from significant limitations and single Gaussian may not be enough.
- **A single Gaussian distribution** fails to capture the two clumps in the data and indeed places much of its probability mass in the central region between the clumps where the data are relatively sparse.
- **A linear combination of two Gaussians** gives a better representation of the data.



The blue curves show contours of constant probability density.  
Figure 2.21 in PRML.

# Motivation of Mixture of Gaussians

- Mixture of Gaussian distributions (Linear combination of Gaussians) can represent very **complex densities**.



Example of a Gaussian mixture in one dimension showing 3 Gaussians (in blue) and their sum in red



# Mixture of Gaussians

- We therefore consider a **superposition** of  $K$  Gaussian densities of the following form:

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$

- Each Gaussian density  $\mathcal{N}(x|\mu_k, \Sigma_k)$  is called a **component of the mixture** and has its own mean  $\mu_k$  and covariance  $\Sigma_k$
- **Mixing coefficients**:  $0 \leq \pi_k \leq 1, \sum_{k=1}^K \pi_k = 1$

# Mixture of Gaussians

- Given  $N$  data points, our goal is to estimate the parameters that maximizes the likelihood function.

- Maximize the log of the likelihood function:

$$\ln p(X|\pi, \mu, \Sigma) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k) \right\}$$

- Difficulty:
  - The likelihood is highly non-linear
  - No closed-form solution

# Chicken-and-Egg Problem

- **Problem:**

- If we know an estimation for the parameters, then we can estimate the cluster of  $x$ .
- If we know the cluster that the data point belongs to, then we can estimate the parameters.
- We don't know either of the above.

- **Solution:**

- Introducing **latent variables** about the cluster assignments.

# Mixture of Gaussians: Introducing Latent Variables

- We introduce a  **$K$ -dimensional binary random variable  $z$**  having a 1-of- $K$  representation in which a particular element  $z_k = 1$  and all other elements are equal to 0.
- Thus,  $z_k \in \{0,1\}$ ,  $\sum_k z_k = 1$ .
- The marginal distribution over  $z$  is specified in terms of the mixing coefficients:  
$$p(z_k = 1) = \pi_k$$
- Because  $z$  uses a 1-of- $K$  representation, we can write the distribution in this form:

$$p(z) = \prod_{k=1}^K \pi_k^{z_k}$$

# Mixtures of Gaussians: Conditional Distribution

- Similarly, the conditional distribution of  $x$  given a particular value for  $z$  is a Gaussian:

$$p(x|z_k = 1) = \mathcal{N}(x|\mu_k, \Sigma_k)$$

which can be written as:

$$p(x|z) = \prod_{k=1}^K \mathcal{N}(x|\mu_k, \Sigma_k)^{z_k}$$

- The marginal distribution of  $x$  is then obtained by summing the joint distribution over all possible values of  $z$ :

$$p(x) = \sum_z p(z)p(x|z) = \sum_{k=1}^K \pi_k \mathcal{N}(x|\mu_k, \Sigma_k)$$

# Mixtures of Gaussians

- The conditional probability of  $z$  given  $x_n$  can be obtained using Bayes' theorem:

$$\gamma(z_{nk}) = p(z_k = 1|x_n) = \frac{p(z_k = 1)p(x_n|z_k = 1)}{\sum_j p(z_j = 1)p(x_n|z_j = 1)} = \frac{\pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)}$$

- We consider  $\pi_k$  as the prior probability of  $z_k = 1$ .
- The quantity  $\gamma(z_{nk})$  will be the corresponding posterior probability once we have observed  $x_n$ .

# Maximum Likelihood

- Gaussian Mixture Distribution:

$$p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x | \mu_k, \Sigma_k)$$

- **The parameters:**  $\pi, \mu, \Sigma$
- Suppose we have a data set of observations  $X = \{x_1, x_2, \dots, x_N\} \in \mathcal{R}^{N \times D}$ , model the data using a mixture of Gaussians.
- **Maximize the log of the likelihood function:**

$$\ln p(X | \pi, \mu, \Sigma) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k) \right\}$$

# EM for Mixtures of Gaussians

- An elegant and powerful method for finding maximum likelihood solutions for Gaussian Mixture Model is called the **expectation-maximization algorithm (EM)**.
- Find  $\pi_k$ ,  $\mu_k$  and  $\Sigma_k$ .
- Consider the **gradients** of log-likelihood function w.r.t these parameters.

$$\ln p(X|\pi, \mu, \Sigma) = \sum_{n=1}^N \ln \left\{ \sum_{k=1}^K \pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k) \right\}$$

$$\mathcal{N}(x|\mu, \Sigma) = \frac{1}{(2\pi)^{D/2}} \frac{1}{|\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$



# EM for Mixtures of Gaussians : $\mu_k$

- Setting the derivatives of  $\ln p(X|\pi, \mu, \Sigma)$  w.r.t  $\mu_k$  to 0:

$$0 = - \sum_{n=1}^N \underbrace{\frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}}_{r(z_{nk})} \Sigma_k (x_n - \mu_k)$$

- We get:

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r(z_{nk}) x_n$$

where  $N_k = \sum_{n=1}^N r(z_{nk})$

# EM for Mixtures of Gaussians: $\mu_k$

- We have:

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r(z_{nk}) x_n$$
$$N_k = \sum_{n=1}^N r(z_{nk})$$

- $r(z_{nk})$ : posterior probability  $p(z_k = 1 | x_n)$ , responsibility that component  $k$  takes for “explaining” the observation  $x_n$ .
- $N_k$  as the **effective number** of points assigned to cluster  $k$ .
- $\mu_k$  for the  $k$ -th Gaussian component is obtained by taking a weighted mean of **ALL of the points** in the data set, in which the weighting factor for data point  $x_n$  is given by the posterior probability  $r(z_{nk})$  that component  $k$  was responsible for generating  $x_n$ .

# EM for Mixtures of Gaussians: $\Sigma_k$

- Setting the derivatives of  $\ln p(X|\pi, \mu, \Sigma)$  w.r.t  $\Sigma_k$  to 0:
- We get:

$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N r(z_{nk})(x_n - \mu_k)(x_n - \mu_k)^T$$
$$N_k = \sum_{n=1}^N r(z_{nk})$$

- Each data point weighted by the corresponding posterior probability.  
Denominator given by the effective number of points associated with the corresponding component.

# EM for Mixtures of Gaussians: $\pi_k$

- Finally, we maximize  $\ln p(X|\pi, \mu, \Sigma)$  w.r.t  $\pi_k$ :
- We know that  $\sum_k \pi_k = 1$ :
- Using a Lagrange multiplier and maximizing the following quantity:

$$\ln p(X|\pi, \mu, \Sigma) + \lambda(\sum_{k=1}^K \pi_k - 1)$$

- It gives:

$$0 = \sum_{n=1}^N \frac{\mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)} + \lambda$$

- Multiply both sides by  $\pi_k$  and sum over  $k$ , make use of  $\sum_k \pi_k = 1$ , we find  $\lambda = -N$  and we get:

$$\pi_k = \frac{N_k}{N} = \frac{\sum_{n=1}^N r(z_{nk})}{N}$$

# EM Algorithm: Initialization

- Initialize the means  $\mu_k$ , covariances  $\Sigma_k$  and mixing coefficient  $\pi_k$ , and evaluate the initial value of the log likelihood.

# EM algorithm: E step

- **E step.** Evaluate the responsibilities using the current parameter values:


$$\gamma(z_{nk}) = \frac{\pi_k \mathcal{N}(x_n | \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n | \mu_j, \Sigma_j)}$$

# EM algorithm: M Step

- **M step.** Re-estimate the parameters using the current responsibilities (where  $N_k = \sum_{n=1}^N r(z_{nk})$ )

$$\mu_k = \frac{1}{N_k} \sum_{n=1}^N r(z_{nk}) x_n$$
$$\Sigma_k = \frac{1}{N_k} \sum_{n=1}^N r(z_{nk}) (x_n - \mu_k)(x_n - \mu_k)^T$$
$$\pi_k = \frac{N_k}{N}$$


# EM for GMM: Example



Initialization: initialize the means, covariance and mixing coefficient [C. Bishop, PRML]




# EM for GMM: Example




E step: use the current values for the parameters to evaluate the posterior probabilities, or responsibilities  $r(z_{nk})$ . [C. Bishop, PRML]

# EM for GMM: Example




M step: use the posterior  $r(z_{nk})$  to re-estimate the means, covariances, and mixing coefficients [C. Bishop, PRML]

# EM for GMM: Example




After 2 complete cycle of EM. [C. Bishop, PRML]

# EM for GMM: Example



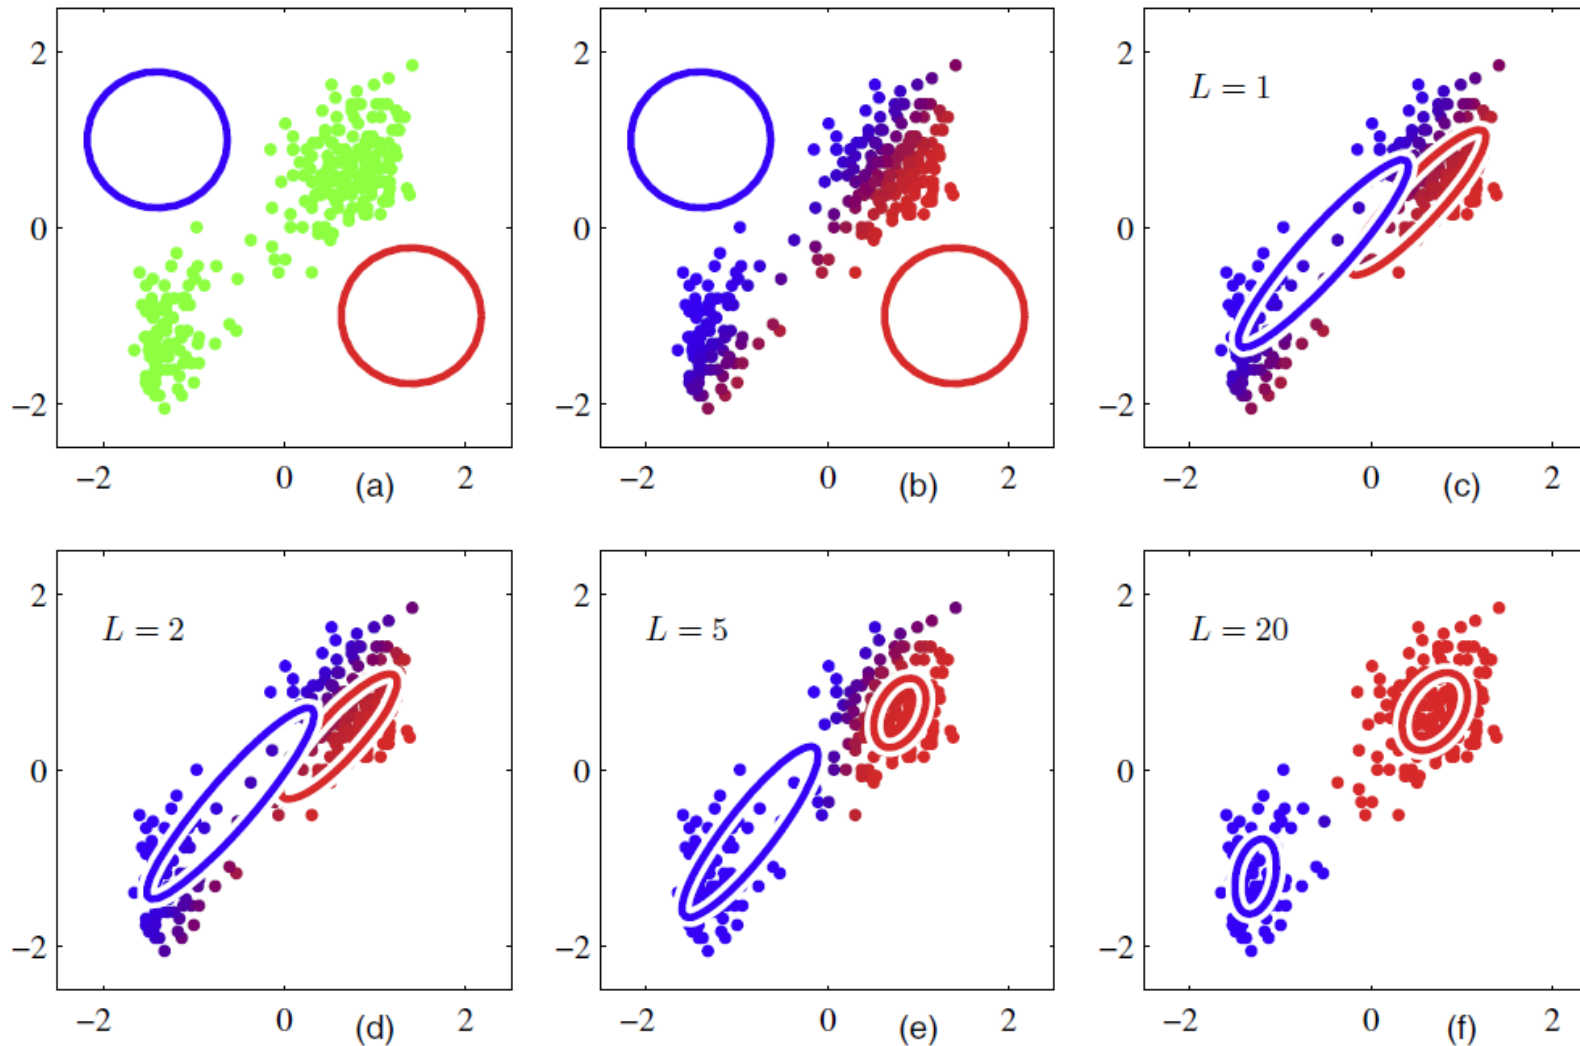
After 5 complete cycle of EM. [C. Bishop, PRML]

# EM for GMM: Example

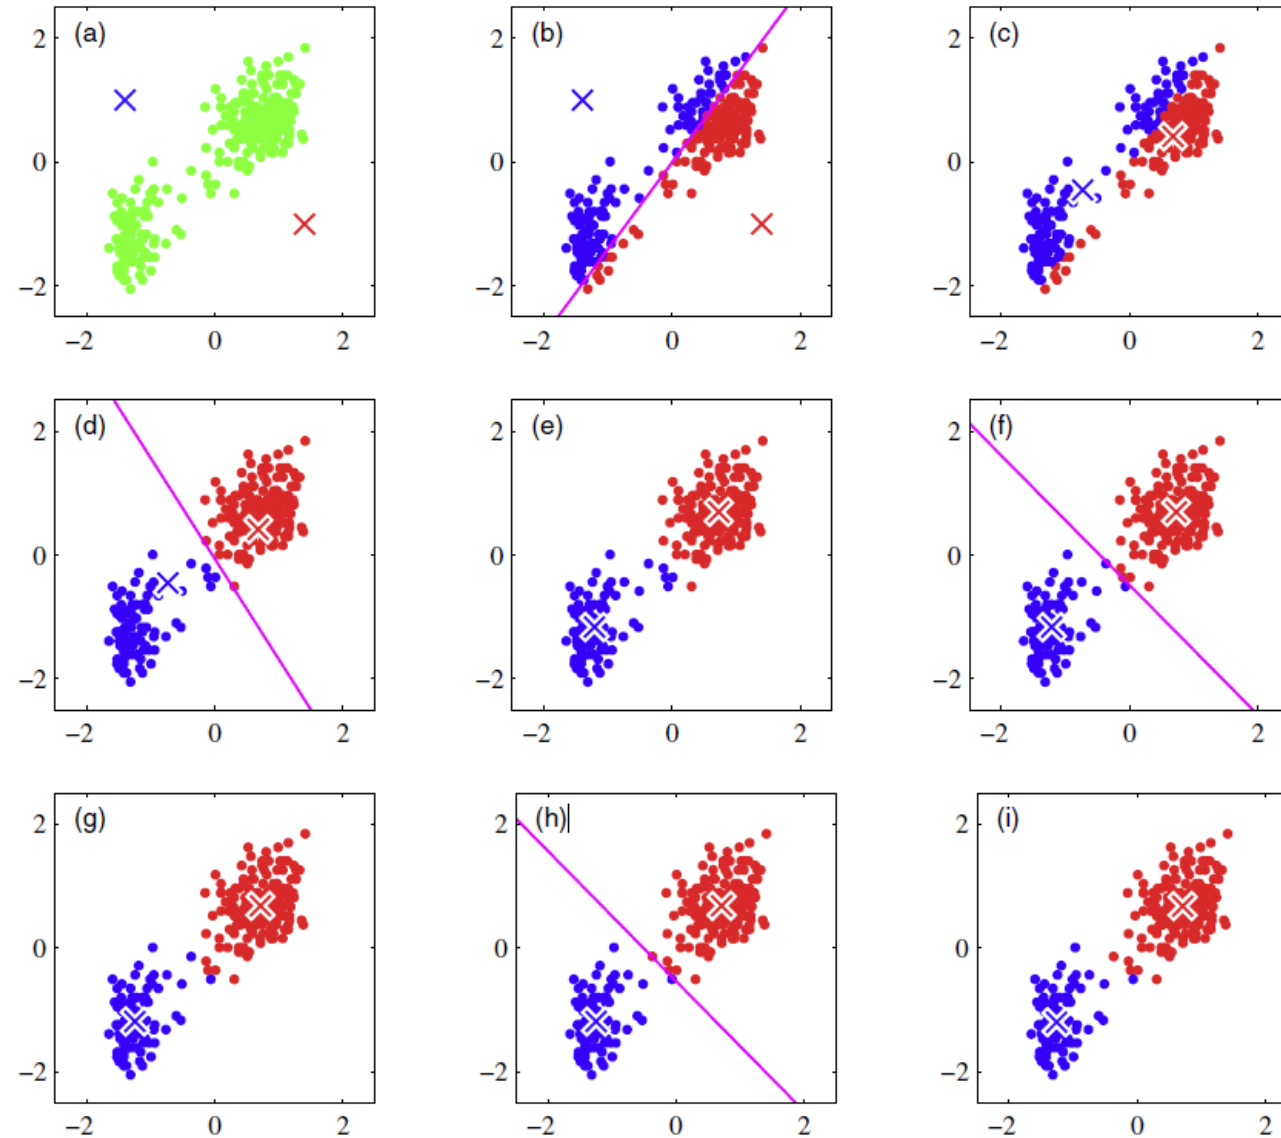


After 20 complete cycle of EM. [C. Bishop, PRML]

# EM for GMM: Soft Assignment



# K-means: Hard Assignment



# K-means vs GMM

- K-means: hard assignment, each data point is associated uniquely with one cluster.
- GMM: soft assignment, based on the posterior probabilities.
  - K-means as special case: consider same and infinitely small variance for each Gaussian component.
  - K-means results can be used as the initialization for EM algorithm.



# Summary of Today's Lecture

- Clustering
- K-means Algorithm
- Gaussian Mixture Model (GMM)